

The Global Arrays User's Manual

(Pacific Northwest National Laboratory Technical Report Number PNNL-13130)

Jarek Nieplocha, Jialin Ju, Manoj Kumar Krishnan, Bruce Palmer, Vinod Tipparaju

This document is intended to be used with the version 3.2 of Global Arrays

October 1, 2002

DISCLAIMER

This material was prepared as an account of work sponsored by an agency of the United States Government. Neither the United States Government nor the United States Department of Energy, nor Battelle, nor any of their employees, MAKES ANY WARRANTY, EXPRESS OR IMPLIED, OR ASSUMES ANY LEGAL LIABILITY OR RESPONSIBILITY FOR THE ACCURACY, COMPLETENESS, OR USEFULNESS OF ANY INFORMATION, APPARATUS, PRODUCT, SOFTWARE, OR PROCESS DISCLOSED, OR REPRESENTS THAT ITS USE WOULD NOT INFRINGE PRIVATELY OWNED RIGHTS.

ACKNOWLEDGMENT

This software and its documentation were produced with United States Government support under Contract Number DE-AC06-76RLO-1830 awarded by the United States Department of Energy. The United States Government retains a paid-up non-exclusive, irrevocable worldwide license to reproduce, prepare derivative works, perform publicly and display publicly by or for the US Government, including the right to distribute to other US Government contractors.

Contents

1. Introduction

- 1.1 Overview
- 1.2 Basic Functionality
- 1.3 Programming Model
- 1.4 Application Guidelines

2. Writing, Building and Running GA Programs

- 2.1 Platform and Library Dependencies
 - 2.1.1 Supported Platforms
 - 2.1.2 Selection of the Communication Network for ARMCI
 - 2.1.3 Selection of the Message passing Library
 - 2.1.4 Dependencies on Other Software
- 2.2 Writing GA Programs
- 2.3 Building GA Programs
 - 2.3.1 Unix Environment
 - 2.3.2 Windows Environment
 - 2.3.3 Writing and building new GA programs
- 2.4 Running GA Programs

3. Initialization and Termination of GA Programs

- 3.1 Message Passing
- 3.2 Memory Allocation
 - 3.2.1 How to determine what the values of MA stack and heap size should be?
- 3.3 GA initialization

- 3.3.1 Limiting Memory Usage by Global Arrays
 - 3.4 Termination
 - 3.5 Creating Arrays
 - 3.5.1 Creating Arrays with Ghost Cells ^{NEW!}
 - 3.6 Destroying Array
 - 4. **One-sided Operations**
 - 4.1 Put/Get
 - 4.2 Acc/Read_inc
 - 4.3 Scatter/Gather
 - 4.4 Periodic Interfaces
 - 5. **Interprocess Synchronization**
 - 5.1 Lock and Mutex
 - 5.2 Fence
 - 5.3 Sync
 - 6. **Collective Array operations**
 - 6.1 Basic Array Operations
 - 6.1.1 Whole Arrays
 - 6.1.2 Patches
 - 6.2 Linear Algebra
 - 6.2.1 Whole Arrays
 - 6.2.2 Patches
 - 6.2.3 Element-wise operations ^{NEW!}
 - 6.3 Interfaces to Third Party Software Packages
 - 6.3.1 Scalapack
 - 6.3.2 PeIGS
 - 6.3.3 Interoperability with Others
 - 6.4 Synchronization Control in Collective Operations ^{NEW!}
 - 7. **Utility Operations**
 - 7.1 Locality Information
 - 7.1.1 Process Information
 - 7.1.2 Cluster Information ^{NEW!}
 - 7.2 Memory Availability
 - 7.3 Message-Passing Wrappers to Reduce/Broadcast Operations
 - 7.4 Others
 - 7.4.1 Inquire
 - 7.4.2 Print
 - 7.4.3 Miscellaneous
 - 8. **GA++: C++ Bindings for Global Arrays** ^{NEW!}
 - 8.1 Overview
 - 8.2 GA++ Classes
 - 8.3 Initialization and Termination
 - 8.4 GAservices
 - 8.5 GlobalArray

Appendix

- A. List of C functions

B. List of Fortran functions

1. Introduction

1.1 Overview

The Global Arrays (GA) toolkit provides a shared memory style programming environment in the context of distributed array data structures (called "global arrays"). From the user perspective, a global array can be used as if it was stored in shared memory. All details of the data distribution, addressing, and data access are encapsulated in the global array objects. Information about the actual data distribution and locality can be easily obtained and taken advantage of whenever data locality is important. The primary target architectures for which GA was developed are massively-parallel distributed-memory and scalable shared-memory systems.

GA divides logically shared data structures into "local" and "remote" portions. It recognizes variable data transfer costs required to access the data depending on the proximity attributes. A local portion of the shared memory is assumed to be faster to access and the remainder (remote portion) is considered slower to access. These differences do not hinder the ease-of-use since the library provides uniform access mechanisms for all the shared data regardless where the referenced data is located. In addition, any processes can access a local portion of the shared data directly/in-place like any other data in process local memory. Access to other portions of the shared data must be done through the GA library calls.

GA was designed to complement rather than substitute the message-passing model, and it allows the user to combine shared-memory and message-passing styles of programming in the same program. GA inherits an execution environment from a message-passing library (w.r.t. processes, file descriptors etc.) that started the parallel program.

GA is implemented as a library with C and Fortran-77 bindings, and there have been also a Python and C++ interfaces (included starting with the release 3.2) developed. Therefore, explicit library calls are required to use the GA model in a parallel C/Fortran program.

A disk extension of the Global Array library is supported by its companion library called Disk Resident Arrays (DRA). DRA maintains array objects in secondary storage and allows transfer of data to/from global arrays.

1.2 Basic Functionality

The basic shared memory operations supported include *get*, *put*, *scatter* and *gather*. They are complemented by **atomic** *read-and-increment*, *accumulate* (reduction operation that combines data in local memory with data in the shared memory location), and *lock* operations. However, these operations can only be used to access data in global arrays rather than arbitrary memory locations. At least one global array has to be created before data transfer operations can be used. These GA operations are truly one-sided/unilateral and will complete regardless of actions taken by the remote process(es) that own(s) the referenced data. In particular, GA does not offer or rely on a polling operation or require inserting any other GA library calls to assure communication progress on the remote side.

A programmer in the GA program has a full control over the distribution of global arrays. Both regular and irregular distributions are supported, see Section 3 for details.

The GA data transfer operations use an array index-based interface rather than addresses of the shared data. Unlike other systems based on global address space that support remote memory (*put/get*) operations, GA

does not require the user to specify the target process/es where the referenced shared data resides — it simply provides a global view of the data structures. The higher level array oriented API (application programming interface) makes GA easier to use, at the same time without compromising data locality control. The library internally performs global array index-to-address translation and then transfers data between appropriate processes. If necessary, the programmer is always able to inquire:


- where and an element or array section is located, and
- which process or processes own data in the specified array section.

The GA toolkit supports four data types in Fortran: integer, real, double precision, and double complex. In the C interface, int, long, float, double and struct double complex are available. Underneath, the library represents the data using C datatypes. For the Fortran users, it means that some arrays created in C for which there is no appropriate datatype mapping to Fortran (for example on the Cray T3E Fortran real is not implemented whereas C float is) might not be accessible. In all the other cases, the datatype representation is transparent.

The supported array dimensions range from one to seven. This limit follows the Fortran convention. The library can be reconfigured to support more than 7-dimensions but only through the C interface.

1.3 Programming Model

The Global Arrays library supports two programming styles: task-parallel and data-parallel. The GA task-parallel model of computations is based on the explicit remote memory copy: The remote portion of shared data has to be copied into the local memory area of a process before it can be used in computations by that process. Of course, the "local" portion of shared data can always be accessed directly thus avoiding the memory copy.



The data distribution and locality control are provided to the programmer. The data locality information for the shared data is also available. The library offers a set of operations for management of its data structures, one-sided data transfer operations, and supportive operations for data locality control and queries. The GA shared memory consistency model is a result of a compromise between the ease of use and a portable performance. The load and store operations are guaranteed to be **ordered** with respect to each other only if they target overlapping memory locations. The store operations (*put*, *scatter*) and *accumulate* complete locally

before returning i.e., the data in the user local buffer has been copied out but not necessarily completed at the remote side. The memory consistency is only guaranteed for:

- multiple read operations (as the data does not change),
- multiple accumulate operations (as addition is commutative), and
- multiple disjoint put operations (as there is only one writer for each element).

The application can manage consistency of its data structures in other cases by using *lock*, *barrier*, and *fence* operations available in the library.

The data-parallel model is supported by a set of collective functions that operate on global arrays or their portions. Underneath, if any interprocessor communication is required, the library uses remote memory copy (most often) or collective message-passing operations.

1.4 Application Guidelines

These are some guidelines regarding suitability of the GA for different types of applications.

When to use GA:

Algorithmic Considerations

- ◇ applications with dynamic and irregular communication patterns
- ◇ for calculations driven by dynamic load balancing
- ◇ need 1-sided access to shared data structures
- ◇ need high-level operations on distributed arrays and/or for out-of-core array-based algorithms (GA + DRA)

Usability Considerations

- ◇ data locality must be explicitly available
- ◇ when coding in message passing becomes too complicated
- ◇ when portable performance is important
- ◇ need object orientation without the overhead of C++

When not to use GA:

Algorithmic Considerations

- ◆ for systolic, or nearest neighbor communications with regular communication patterns
- ◆ when synchronization associated with cooperative point-to-point message passing is needed (e.g., Cholesky factorization in Scalapack)

Usability Considerations

- ◆ when interprocedural analysis and compiler parallelization is more effective
- ◆ a parallel language support is sufficient and robust compilers available

2. Writing, Building and Running GA Programs

The web page www.emsl.pnl.gov:2080/docs/global/support.html contains updated information about using GA on different platforms. Please refer to this page frequently for most recent updates and platform information.

2.1 Platform and Library Dependencies

2.1.1 Supported Platforms

- IBM SP, CRAY T3E/J90/SV1, SGI Origin, Fujitsu VX/VPP, Hitachi
- Cluster of workstations: Solaris, IRIX, AIX, HPUX, Digital/Tru64 Unix, Linux, NT
- Standalone uni- or multi-processor workstations or servers
- Standalone uni- or multi-processor Windows NT workstations or servers

Older versions of GA supported some additional (now obsolete) platforms such as: IPSC, KSR, PARAGON, DELTA, CONVEX. They are not supported in the newer (>3.1) versions because we do not have access to these systems. We recommend using GA 2.4 on these platforms.

For most of the platforms, there are two versions available: 32-bit and 64-bit.

Platform	32-bit TARGET name	64-bit TARGET name	Remarks
Sun ultra	SOLARIS	SOLARIS64	64-bit version added in GA 3.1
IBM RS/6000	IBM	IBM64	64-bit version added in GA 3.1
IBM SP	LAPI	not available	no support yet for user-space communication in the 64-bit mode by IBM
Compaq/DEC alpha	not available	DECOSF	
HP pa-risc	HPUX	HPUX64	64-bit version added in GA 3.1
Linux x86, ultra, powerpc	LINUX	not available	
Linux IA64 (Itanium)	not available	LINUX64	
Linux alpha	not available	LINUX64	64-bit version added in GA 3.1; Compaq compilers rather than GNU required
Cray T3E	not available	CRAY-T3E	
Cray J90	not available	CRAY-YMP	
Cray SV1	not available	CRAY-SV1	
SGI IRIX mips	SGI_N32, SGI	SGITFP	
Hitachi SR8000	HITACHI	not available	

To aid development of fully portable applications, in 64-bit mode Fortran integer datatype is 64-bits. It is motivated by 1) the need of applications to use very large data structures and 2) Fortran INTEGER*8 not being fully portable. The 64-bit representation of integer datatype is accomplished by using appropriate Fortran compiler flag.

Because of limited interest in heterogenous computing among known us GA users, the Global Array library *still does not support heterogeonous platforms*. This capability can be added if required by new applications.

2.1.2 Selection of the communication network for ARMCI

Some cluster installations can be equipped with a high performance network which offer instead, or in addition to TCP/IP some special communication protocol, for example GM on Myrinet network. To achieve high performance in Global Arrays, ARMCI must be built to use these protocols in its implementation of one-sided communication. Starting with GA 3.1, this is accomplished by setting an environment variable `ARMCI_NETWORK` to specify the protocol to be used. In addition, the it might be necessary to provide location for the header files and library path corresponding to location of s/w supporting the appropriate protocol API, see `g/armci/config/makecoms.h` for details.

Network	Protocol name	ARMCI_NETWORK setting	Supported platforms
Ethernet	TCP/IP	SOCKETS (optional/default)	workstation clusters
Quadrics	Elan/Shmem	QUADRICS	Linux (alpha,x86,IA64), Compaq
Myrinet	GM	GM	Linux (x86,ultra,IA64)
Giganet cLAN	VIA	VIA	Linux (x86)

The port on top of Myrinet has been partially optimized. The Giganet/VIA port has not been optimized yet and is included on the experimental basis.

2.1.3 Selection of the message-passing library

As explained in Section 3, GA works with either MPI or TCGMSG message-passing libraries. That means that GA applications can use either of these interfaces. Selection of the message-passing library takes place when GA is built. Since the TCGMSG library is small and compiles fast, it is included with the GA distribution package and built on Unix workstations by default so that the package can be built as fast and as conveniently to the user as possible. There are three possible configurations for running GA with the message-passing libraries:

1. with TCGMSG
2. with MPI and TCGMSG emulation library: TCGMSG-MPI, that implements functionality of TCGMSG using MPI. In this mode, the message passing library is initialized using a TCGMSG `PBEGIN(F)` call which internally references `MPI_Initialize`. To enable this mode, define the environmental variable `USE_MPI`.
3. directly with MPI. In this mode, GA program should contain MPI initialization calls instead of `PBEGIN(F)`.

For the MPI versions, the optional environmental variables `MPI_LIB` and `MPI_INCLUDE` are used to point to the location of the MPI library and include directories if they are not in the standard system location(s). GA programs are started with the mechanism that any other MPI programs use on the given platform.

The recent versions of MPICH (an MPI implementation from ANL/Mississippi State) keep the MPI header files in more than one directory and provide compiler wrappers that implicitly point to the appropriate header files. One can :

- use `MPI_INCLUDE` by expanding the string with another directory component prefixed with "-I" (you are passing include directory names as a part of compiler flags), or (starting with GA 3.1) separated by comma "," and without the prefix, OR
- use MPI aware compiler wrappers e.g., `mpicc` and `mpif77` to build GA right out of the box on UNIX workstations:

`make FC=mpif77 CC=mpicc` One disadvantage of the second approach is that GA makefile in some circumstances might be not able to determine which compiler (e.g., GNU or PGI) is called underneath by the MPICH compiler wrappers. Since different compilers provide different Fortran/C interface, the package might fail to build. This problem is most likely to occur on non-Linux Unix systems with non-native compilers (e.g., gcc).

On Windows NT, the current version of GA was tested with WMPI, an NT implementation derived from MPICH in Portugal.

2.1.3 Dependencies on other software

In addition to the message-passing library, GA requires:

- MA (Memory Allocator), a library for management of local memory;
- ARMCI, a one-sided communication library that GA uses as its run-time system;
- BLAS library is required for the eigensolver and `ga_dgemm`;
- LAPACK library is required for the eigensolver (a subset is included with GA, which is built into *liblinalg.a*);

GA may also depend on other software depending on the functions being used.

- GA *eigensolver*, `ga_diag`, is a wrapper for the eigensolver from the PEIGS library; (Please contact [George Fann <gi_fann@pnl.gov>](mailto:gi_fann@pnl.gov) about PEIGS)
- SCALAPACK, PBLAS, and BLACS libraries are required for `ga_lu_solve`, `ga_cholesky`, `ga_llt_solve`, `ga_spd_invert`, `ga_solve`. If these libraries are not installed, the named operations will not be available.
- If one would like to generate trace information for GA calls, an additional library *libtrace.a* is required, and the `-DGA_TRACE` define flag should be specified for C and Fortran compilers.

2.2 Writing GA Programs

C programs that use Global Arrays should include files `'global.h'`, `'ga.h'`, `'macdecls.h'`. Fortran programs should include the files `'mafdecls.fh'`, `'global.fh'`. Fortran source must be preprocessed as a part of compilation.

The GA program should look like:

- When GA runs with MPI

Fortran**C**

```

call mpi_init(..)    MPI_Init(..)      ! start MPI
call ga_initialize()  GA_Initialize()    ! start global arrays
status = ma_init(..) MA_Init(..)       ! start memory allocator

.... do work          .... do work

call ga_terminate()  GA_Terminate()     ! tidy up global arrays
call mpi_finalize()  MPI_Finalize()     ! tidy up MPI
stop                  ! exit program

```

- When GA runs with TCGMSG or TCGMSG-MPI

Fortran**C**

```

call pbeginf()       PBEGIN(..)        ! start TCGMSG
call ga_initialize()  GA_Initialize()    ! start global arrays
status = ma_init(..) MA_Init(..)       ! start memory allocator

.... do work          .... do work

call ga_terminate()  GA_Terminate()     ! tidy up global arrays
call pend()          PEND_()            ! tidy up tcgmsg
stop                  ! exit program

```

The *ma_init* call looks like :

```
status = ma_init(type, stack_size, heap_size)
```

and it basically just goes to the OS and gets *stack_size+heap_size* elements of size type. The amount of memory MA allocates need to be sufficient for storing global arrays on some platforms. Please refer to section [3.3.1](#) for the details and information on more advanced usage of MA in GA programs.

2.3 Building GA Programs

Use *GNU make* to build the GA library and application programs on Unix and Microsoft *nmake* on Windows. The structure of the available makefiles are

- GNUmakefile: Unix makefile
- MakeFile: Windows NT makefile
- Makefile.h: definitions & include symbols

The user needs to specify TARGET in the GNUmakefile or on the command line when calling make. The library and test programs should be built by calling make in the current directory. Valid TARGETs are listed by calling make in the top level distribution directory on UNIX family of systems when TARGET is not defined. On Windows, WIN32, CYGNUS and INTERIX (previously known as OpenNT) are supported.

One could affect which compilers and compiler flags the package uses (instead of the predefined defaults) by specifying them for GNU make on the command line;

- CC – name of the C compiler (e.g., gcc, cc, or ccc)
- FC – name of the Fortran compiler (e.g., g77, f90, mpif77 or fort)
- COPT – optimization or debug flags for the C compiler (e.g., -g, -O3)
- FOPT – optimization or debug flags for the Fortran compiler (e.g., -g, -O1)

For example,

```
gmake FC=f90 CC=ccc FOPT=-O4 COPT=-g
```

Note that GA provides only Fortran-77 interfaces. To use and compile with a Fortran 90 compiler, it has to support a subset of Fortran-77.

2.3.1 Unix Environment

To build GA with the MPI, user needs to define environmental variables *USE_MPI*, *MPI_LIB* and *MPI_INCLUDE* which should point to the location of the MPI library and include directories.

Example: using csh/tcsh (assume using MPICH installed in /usr/local on IBM workstation)

```
setenv USE_MPI y
setenv MPI_LOC /usr/local/mpich
setenv MPI_LIB $MPI_LOC/lib/rs6000/ch_shmem
setenv MPI_INCLUDE $MPI_LOC/include
```

Additionally, if the TCGMSG-MPI library is not needed, the make/environmental variable MSG_COMMS should be defined as MSG_COMMS = MPI.

Interface routines to ScaLAPACK are only available with MPI, and of course with ScaLAPACK. The user is required to define the environment variables *USE_SCALAPACK*, and the location of ScaLAPACK & Co. libraries in variable SCALAPACK.

Example: using csh/tcsh

Global Arrays User Manual

```
setenv USE_SCALAPACK y
setenv SCALAPACK '-L/msrc/proj/scalapack/LIB/rs6000
               -lscalapack -lpblas -ltools -lblacsF77cinit -lblacs'
setenv USE_MPI y
```

Since there are certain interdependencies between blacs and blacsF77cinit, some system might require specification of `-lblacs` twice to fix the unresolved external symbols from these libs.

To build the library, type
`make` or `gmake`

To build an application based on GA located in `g/global/testing`, for example, the application's name is `app.c` (or `app.F`, `app.f`), type
`make app.x` or `gmake app.x`

Please refer to compiler flags in file `g/global/Makefile.h` to make sure that Fortran and C compiler flags are consistent with flags use to compile your application. This may be critical when Fortran compiler flags are used to change the default length of the integer datatype.

2.3.2 Windows NT

To build GA on Windows NT, MS Power Fortran 4 or DEC Visual Fortran 5 or later, and MS Visual C 4 or later are needed. Other compilers might need the default compilation flags modified. When commercial Windows compilers are not available, one can choose to use CYGNUS or INTERIX and build it as any other Unix box using GNU compilers.

First of all, one needs to set environment variables (same as in Unix environment). GA needs to know where find the MPI include files and libraries. To do this, select the *Environment* tab under the Control Panel, then set the variables to point to the location of MPI, for example for WMPI on disk D:

```
set MPI_INCLUDE as d:\Wmpi\Include
set MPI_LIB as d:\Wmpi\Console
```

Make sure that the dynamic link libraries required by the particular implementation of MPI are copied to the appropriate location for the system DLLs. For WMPI, copy `VWMPI.dll` to `\winnt`.

In the top directory do,

```
nmake
```

The GA `test.exe` program can be built in the `g\global\testing` directory:

```
nmake test.exe
```

In addition, the HPVM package from UCSD offers the GA interface in the NT/Myrinet cluster environment.

GA could be built on Windows 95/98. However, due to the DOS shell limitations, the top level NTmakefile will not work. Therefore, each library has to be made separately in its own directory. The environment variables referring to MPI can be hardcoded in the NT makefiles.

2.3.3 Writing and building new GA programs

2.3.2 Windows NT

For small programs contained in a single file, the most convenient approach is to put your program file into the `g/global/testing` directory. [The existing GNU make suffix rules would build an executable with the ".x" suffix from any C or Fortran source file.](#) You do not have to modify makefiles in `g/global/testing` at all. For example, if your program is contained in `myfile.c` or `myfile.F` and you place it in that directory, all you need to do to create an executable called `myfile.x` is to type: `make myfile.x`.

Windows `nmake` is not as powerful as GNU `make` – you would need to modify the NT makefile.

This approach obviously is not feasible for large packages that contain multiple source files and directories. In that case you need to provide appropriate definitions in your makefile:

- to header files located in the include directory, `g/include`, where all public header files are copied in the process of building GA
- add references to `libglobal.a` (Unix) `global.lib` (Windows) and `libma.a` (Unix) `ma.lib` (Windows) in `g/lib/$(TARGET)` and for the message-passing libraries
- follow compilation flags for the GA test programs in GNU and Windows makefiles `g/config/makefile.h`. The recommended approach is to include `g/config/makefile.h` in your makefile.

Starting with GA 3.1, one could simplify linking of applications by including `g/armci/config/makecoms.h` and `g/armci/config/makemp.h` that define all the necessary platform specific libraries that are required by GA.

2.4 Running GA Programs

Assume the `app.x` had already been built. To run it,

1. On MPPs, such as Cray T3E, or IMB SP

Use appropriate system command to specify the number of processors, load and run the programs.

Example: to run on four processors on the Cray T3E, use

```
mpprun -n 4 app.x
```

2. On shared memory systems and (network of) workstations (including linux cluster)

If the `app.x` is built based on MPI, run the program the same way as any other MPI programs.

Example: to run on four processes on SGI workstation, use

```
mpirun -np 4 app.x, or  
app.x -np 4
```

If `app.x` is built based on TCGMSG(not including, Fujitsu, Cray J90, and Windows, because there are no native ports of TCGMSG), to execute the program on Unix workstations/servers, one should use

Global Arrays User Manual

the 'parallel' program (built in tcgmsg/ipcv4.0). After building the application, a file called 'app.x.p' would also be generated (If there is not such a file, make it: `make app.x.p`). This file can be edited to specify how many processors and tasks to use, and how to load the executables. Make sure that 'parallel' is accessible (you might copy it into your 'bin' directory). To execute, type:

```
parallel app.x
```

3. On Microsoft NT, there is no support for TCGMSG, which means you can only build your application based on MPI. Run the application program the same way as any other MPI programs. For, WMPI you need to create the .pg file.

Example:

```
R:\nt\g\global\testing> start /b test.exe
```

3. Initialization and Termination

For historical reasons (the 2-dimensional interface was developed first), many operations have two interfaces, one for two dimensional arrays and the other for arbitrary dimensional (one- to seven- dimensional, to be more accurate) arrays. The latter can definitely handle two dimensional arrays as well. The supported data types are *integer*, *double precision*, and *double complex*. Global Arrays provide C and Fortran interfaces in the same (mixed-language) program to the same array objects. The underlying data layout is based on the Fortran convention.

GA programs require message-passing and Memory Allocator (MA) libraries to work. Global Arrays is an extension to the message-passing interface. GA internally does not allocate local memory from the operating system – all dynamically allocated local memory comes from MA. We will describe the details of memory allocation later in this section.

3.1 Message Passing

The first version of Global Arrays was released in 1994 before robust MPI implementations became available. At that time, GA worked only with TCGMSG, a message-passing library that one of the GA authors (Robert Harrison) had developed before. In 1995, support for MPI was added. At the present time, the GA distribution still includes the TCGMSG library for backward compatibility purposes, and because it is small, fast to compile, and provides a minimal message-passing support required by GA programs. The user can enable the MPI-compatible version of GA by defining `USE_MPI` environment variable before compiling the GA toolkit. On systems where vendors provide MPI with interoperable C and Fortran interfaces, there is no advantage in compiling or using TCGMSG.

The GA toolkit needs the following functionality from any message-passing library it runs with:

- initialization and termination of processes in an SPMD (single-program-multiple-data) program,
- synchronization,
- functions that return number of processes and calling process id,
- broadcast,
- reduction operation for integer and double datatypes, and
- a function to abort the running parallel job in case of an error.

The message-passing library has to be initialized before the GA library and terminated after the GA library is terminated.

GA provides two functions `ga_nnodes` and `ga_nodeid` that return the number of processes and the calling process id in a parallel program. Starting with release 3.0, these functions return the same values as their message-passing counterparts. In earlier releases of GA on clusters of workstations, the mapping between GA and message-passing process ids were nontrivial. In these cases, the `ga_list_nodeid` function (now obsolete) was used to describe the actual mapping.

Although message-passing libraries offer their own barrier (global synchronization) function, this operation does not wait for completion of the outstanding GA communication operations. The GA toolkit offers a `ga_sync` operation that can be used for synchronization, and it has the desired effect of waiting for all the outstanding GA operations to complete.

3.2 Memory Allocation

GA uses a very limited amount of statically allocated memory to maintain its data structures and state. Most of the memory is allocated dynamically as needed, primarily to store data in newly allocated global arrays or as temporary buffers internally used in some operations, and deallocated when the operation completes.

There are two flavors of dynamically allocated memory in GA: shared memory and local memory. Shared memory is a special type of memory allocated from the operating system (UNIX and Windows) that can be shared between different user processes (MPI tasks). A process that attaches to a shared memory segment can access it as if it was local memory. All the data in shared memory is directly visible to every process that attaches to that segment. On shared memory systems and clusters of SMP (symmetric multiprocessor) nodes, shared memory is used to store global array data and is allocated by the Global Arrays run-time system called ARMCI. ARMCI uses shared memory to optimize performance and avoid explicit interprocessor communication within a single shared memory system or an SMP node. ARMCI allocates shared memory from the operating system in large segments and then manages memory in each segment in response to the GA allocation and deallocation calls. Each segment can hold data in many small global arrays. ARMCI does not return shared memory segments to the operating system until the program terminates (calls `ga_terminate`).

On systems that do not offer shared-memory capabilities or when a program is executed in a serial mode, GA uses local memory to store data in global arrays.

All of the dynamically allocated local memory in GA comes from its companion library, the Memory Allocator (MA) library. MA allocates and manages local memory using *stack* and *heap* disciplines. Any buffer allocated and deallocated by a GA operation that needs temporary buffer space comes from the MA *stack*. Memory to store data in global arrays comes from *heap*. MA has additional features useful for program debugging such as:

- left and right guards: they are stamps that detect if a memory segment was overwritten by the application,
- named memory segments, and
- memory usage statistics for the entire program.

Explicit use of MA by the application to manage its non-GA local data structures is not necessary but encouraged. Because MA is used implicitly by GA, it has to be initialized before the first global array is allocated. The `MA_init` function requires users to specify memory for *heap* and *stack*. This is because MA:

- allocates from the operating system only one segment equal in size to the sum of *heap* and *stack*,
- manages both allocation schemes using memory coming from opposite ends of the same segment, and
- the boundary between free *stack* and *heap* memory is dynamic.

It is not important what the stack and heap size argument values are as long as the aggregate memory consumption by a program does not exceed their sum at any given time.

3.2.1 How to determine what the values of MA stack and heap size should be?

The answer to this question depends on the run-time environment of the program including the availability of shared memory. A part of GA initialization involves initialization of the ARMCI run-time library. ARMCI dynamically determines if the program can use shared memory based on the architecture type and current configuration of the SMP cluster. For example, on uniprocessor nodes of the IBM SP shared memory is not

used whereas on the SP with SMP nodes it is. This decision is made at run-time. GA reports the information about the type of memory used with the function `ga_uses_ma()`. This function returns false when shared memory is used and true when MA is used.

Based on this information, a programmer who cares about the efficient usage of memory has to consider the amount of memory per single process (MPI task) needed to store data in global arrays to set the `heap_size` argument value in `ma_init`. The amount of stack space depends on the GA operations used by the program (for example `ga_mulmat_patch` or `ga_dgemm` need several MB of buffer space to deliver good performance) but it probably should not be less than 4MB. The stack space is only used when a GA operation is executing and it is returned to MA when it completes.

3.3 GA Initialization

The GA library is initialized after a message-passing library and before MA. It is possible to initialize GA after MA but it is not recommended: GA must first be initialized to determine if it needs shared or MA memory for storing distributed array data. There are two alternative functions to initialize GA:

```
Fortran  subroutine ga_initialize()
C        void GA_Initialize()
C++      void GA::Initialize(int argc, char **argv)
```

and

```
Fortran  subroutine ga_initialize_ltd(limit)
C        void GA_Initialize_ltd(size_t limit)
C++      void GA::Initialize(int argc, char **argv, size_t limit)
```

The first interface allows GA to consume as much memory as the application needs to allocate new arrays. The latter call allows the programmer to establish and enforce a limit within GA on the memory usage.

Note: In GA++, there is an additional functionality as follows:

```
C++      void GA::Initialize(int argc, char *argv[], unsigned long
heapSize, unsigned long stackSize, int type, size_t limit=0)
```

3.3.1 Limiting Memory Usage by Global Arrays

GA offers an optional mechanism that allows a programmer to limit the aggregate memory consumption used by GA for storing Global Array data. These limits apply regardless of the type of memory used for storing global array data. They do not apply to temporary buffer space GA might need to use to execute any particular operation. The limits are given per process (MPI task) in bytes. If the limit is set, GA would not allocate more memory in global arrays that would exceed the specified value – any calls to allocate new arrays that would simply fail (return false). There are two ways to set the limit:

1. at initialization time by calling `ga_initialize_ltd`, or
2. after initialization by calling the function

```
Fortran  subroutine ga_set_memory_limit(limit)
C        void GA_Set_memory_limit(size_t limit)
C++      void GA::GAServices::setMemoryLimit(size_t limit)
```

It is encouraged that the user choose the first option, even though the user can initialize the GA normally and set the memory limit later.

Example: Initialization of MA and setting GA memory limits

```

      call ga_initialize()
      if (ga_uses_ma()) then
        status = ma_init(MT_DBL, stack, heap+global)
      else
        status = ma_init(mt_dbl,stack,heap)
        call ga_set_memory_limit(ma_sizeof(MT_DBL,global,MT_BYTE))
      endif
      if(.not. status) ... !we got an error condition here

```

In this example, depending on the value returned from `ga_uses_ma()`, we either increase the *heap* size argument by the amount of memory for global arrays or set the limit explicitly through `ga_set_memory_limit()`. When GA memory comes from MA we do not need to set this limit through the GA interface since MA enforces its memory limits anyway. In both cases, the maximum amount of memory acquired from the operating system is capped by the value `stack+heap+global`.

3.4 Termination

The normal way to terminate a GA program is to call the function

```

Fortran  subroutine ga_terminate()
C        void GA_Terminate()
C++      void GA::Terminate()

```

The programmer can also abort a running program for example as part of handling a programmatically detected error condition by calling the function

```

Fortran  subroutine ga_error(message, code)
C        void GA_Error(char *message, int code)
C++      void GA::GAServices::error(char *message, int code)

```

3.5 Creating arrays

There are two way to create new arrays:

1. From scratch, for regular distribution, using

```

n-d Fortran  logical function nga_create(type, ndim, dims, array_name,
                                     chunk, g_a)
2-d Fortran  logical function ga_create(type, dim1, dim2, array_name,
                                     chunk1, chunk2, g_a)
C            int NGA_Create(int type, int ndim, int dims[], char
*array_name,
                                     int chunk[])
C++          GA::GlobalArray* GA::GAServices::createGA(int type, int

```

ndim,

```
int dims[], char *array_name, int chunk[])
```

or for regular distribution, using

```
n-d Fortran logical function nga_create irreg(type, ndim, dims,
                                             array_name,
                                             map, nblock, g_a)
2-d Fortran logical function ga_create irreg(type, dim1, dim2,
                                             array_name,
                                             map1, nblock1, map2, nblock2,
                                             g_a)
C          int NGA_Create irreg(int type, int ndim, int dims[],
C++        GA::GlobalArray* GA::GAServices::createGA(int type, int
                                             ndim,
                                             int dims[], char *array_name, int map[],
int block[])
```


2. Based on a template (an existing array) with the function

```
Fortran logical function ga_duplicate(g_a, g_b, array_name)
C       int GA_Duplicate(int g_a, char *array_name)
C++     int GA::GAServices::duplicate(int g_a, char *array_name)- or


C++     GA::GlobalArray* GA::GAServices::createGA(int g_a, char
*array_name)
```

In this case, the new array inherits all the properties such as distribution, datatype and dimensions from the existing array.

With the regular distribution, the programmer can specify block size for none or any dimension. If block size is not specified the library will create a distribution that attempts to assign the same number of elements to each processor (for static load balancing purposes). The actual algorithm used is based on heuristics.



With the irregular distribution, the programmer specifies distribution points for every dimension using *map* array argument. The library creates an array with the overall distribution that is a Cartesian product of distributions for each dimension. A specific example is given in the documentation.



If an array cannot be created, for example due to memory shortages or an enforced memory consumption limit, these calls return failure status. Otherwise an integer handle is returned. This handle represents a global array object in all operations involving that array. This is the only piece of information the programmer needs to store for that array. All the properties of the object (data type, distribution data, name, number of dimensions and values for each dimension) can be obtained from the library based on the handle at any time, see Section 7.4. It is not necessary to keep track of this information explicitly in the application code.

Note that regardless of the distribution type at most one block can be owned/assigned to a processor.

3.5.1 Creating Arrays with Ghost Cells

Individual processors ordinarily only hold the portion of global array data that is represent by the lo and hi index arrays returned by a call to `nga_distribution` or that have been set using the `nga_create_irreg` call. However, it is possible to create global arrays where this data is padded by a boundary region of array elements representing portions of the global array residing on other processors. These boundary regions can be updated with data from neighboring processors by a call to a single GA function. To create global arrays with these extra data elements, referred to in the following as ghost cells, the user needs to call either the functions:

```

n-d Fortran  logical function nga_create_ghosts(type, dims, width,
                                                    array_name,
                                                    chunk, g_a)

C            int nga_create_ghosts(int type, int ndim, int dims[],
                                      int width[],
                                      char *array_name, int chunk[])

C++          int GA::GAServices::createGA_Ghosts(int type, int ndim,
                                                    int dims[],
                                                    int width[], char *array_name, int
                                                    chunk[])

n-d Fortran  logical function nga_create_ghosts_irreg(type, dims,
                                                         width,
                                                         array_name, map, block, g_a)

C            int nga_create_ghosts_irreg(int type, int ndim, int
                                           dims[],
                                           int width[], char *array_name, int map[], int
                                           block[])

C++          int GA::GAServices::createGA_Ghosts(int type, int ndim,
                                                    int dims[],

```


Global Arrays User Manual

```
int width[], char *array_name, int map[], int  
block[])
```

These two functions are almost identical to the `nga_create` and `nga_create_irreg` functions described above. The only difference is the parameter array `width`. This is used to control the width of the ghost cell boundaries in each dimension of the global array. Different dimensions can be padded with different numbers of ghost cells, although it is expected that for most applications the widths will be the same for all dimensions. If the width has been set to zero for all dimensions, then these two functions are completely equivalent to the functions `nga_create` and `nga_create_irreg`.


To illustrate the use of these functions, an ordinary global array is shown below. The boundaries represent the data that is held on each processor.

Global Array



For a global array with ghost cells, the data distribution can be visualized as follows:

Global Array with Ghost Cells



Each processor holds visible data, corresponding to the data held on each processor of an ordinary global array, and ghost cell data, corresponding to neighboring points in the global array that would ordinarily be held on other processors. This data can be updated in a single call to `nga_update`, described under the collective operations section of the user documentation. Note that the ghost cell data duplicates some portion of the data in the visible portion of the global array. The advantage of having the ghost cells is that this data ordinarily resides on other processors and can only be retrieved using additional calls. To access the data in the ghost cells, the user must use the `nga_access_ghosts` function described in Section 6.1.

3.6 Destroying arrays

Global arrays can be destroyed by calling the function

```
Fortran  subroutine ga_destroy(g_a)
```

```
C      void GA_Destroy(int g_a)
C++    void GA::GlobalArray::destroy( )
```

that takes as its argument a handle representing a valid global array. It is a fatal error to call `ga_destroy` with a handle pointing to an invalid array.

All active global arrays are destroyed implicitly when the user calls `sga_terminate`.

4. One-sided Operations

Global Arrays provide one-sided, noncollective communication operations that allow to access data in global arrays without cooperation with the process or processes that hold the referenced data. These processes do not know what data items in their own memory are being accessed or updated by remote processes. Moreover, since the GA interface uses global array indices to reference nonlocal data, the calling process does not even have to know process ids and location in memory where the referenced data resides.

The one-sided operations that Global Arrays provide can be summarized into three categories:

Remote blockwise write/read	<code>ga_put</code> , <code>ga_get</code>
Remote atomic update	<code>ga_acc</code> , <code>ga_read_inc</code> , <code>ga_scatter_acc</code>
Remote elementwise write/read	<code>ga_scatter</code> , <code>ga_gather</code>

4.1 Put/Get

Put and *get* are two powerful operations for interprocess communication, performing remote write and read. Because of their one-sided nature, they don't need cooperation from the process(es) that owns the data. The semantics of these operations do not require the user to specify which remote process or processes own the accessed portion of a global array. The data is simply accessed as if it were in shared memory.

Put copies data from the local array to the global array section, which is

```

n-D Fortran  subroutine nga_put(g_a, lo, hi, buf, ld)
2-D Fortran  subroutine ga_put(g_a, ilo, ihi, jlo, jhi, buf, ld)
C            void _NGA_Put(int g_a, int lo[], int hi[], void *buf, int
               ld[])
C++          void _GA::GlobalArray::put(int lo[], int hi[], void *buf,
               int ld[])

```

All the arguments are provided in one call: `lo` and `hi` specify where the data should go in the global array; `ld` specifies the stride information of the local array `buf`. The local array should have the same number of dimensions as the global array; however, it is really required to present the n-dimensional view of the local memory buffer, that by itself might be one-dimensional.

The operation is transparent to the user, which means the user doesn't have to worry about where the region defined by `lo` and `hi` is located. It can be in the memory of one or many remote processes, owned by the local process, or even mixed (part of it belongs to remote processes and part of it belongs to a local process).

Get is the reverse operation of *put*. It copies data from a global array section to the local array. It is

```

n-D Fortran  subroutine nga_get(g_a, lo, hi, buf, ld)
2-D Fortran  subroutine ga_get(g_a, ilo, ihi, jlo, jhi, buf, ld)
C            void _NGA_Get(int g_a, int lo[], int hi[], void *buf, int
               ld[])
C++          void _GA::GlobalArray::get(int lo[], int hi[], void *buf,

```


```
int ld[])
```

Similar to *put*, *lo* and *hi* specify where the data should come from in the global array, and *ld* specifies the stride information of the local array *buf*. The local array is assumed to have the same number of dimensions as the global array. Users don't need to worry about where the region defined by *lo* and *hi* is physically located.

Example:

For a *ga_get* operation transferring data from the (11:15,1:5) section of a 2-dimensional 15 x10 global array into a local buffer 5 x10 array we have: (In Fortran notation)

lo={11,1}, *hi*={15,5}, *ld*={10}



4.2 Accumulate and read-and-increment

It is often useful in a *put* operation to combine the data moved to the target process with the data that resides at that process, rather than replacing the data there. *Accumulate* and *read_inc* perform **atomic** remote update to a patch (a section of the global array) in the global array and an element in the global array, respectively. They don't need the cooperation of the process(es) who owns the data. Since the operations are atomic, the same portion of a global array can be referenced by these operations issued by multiple processes and the GA will assure the correct and consistent result of the updates.

Accumulate combines the data from the local array with data in the global array section, which is

```
n-D Fortran  subroutine nga_acc(g_a, lo, hi, buf, ld, alpha)
2-D Fortran  subroutine ga_acc(g_a, ilo, ihi, jlo, jhi, buf, ld,
               alpha)
C            void nga_acc(int g_a, int lo[], int hi[], void *buf, int
               ld[],
               void *alpha)
C++          void NGA::GlobalArray::acc(int lo[], int hi[], void *buf,
               int ld[],
               void *alpha)
```

The local array is assumed to have the same number of dimensions as the global array. Users don't need to worry about where the region defined by *lo* and *hi* is physically located. The function performs

*global array section (lo[], hi[]) += alpha * buf*

Read_inc remotely updates a particular element in the global array, which is

```

n-D Fortran  subroutine nga_read_inc(g_a, subscript, inc)
2-D Fortran  subroutine ga_read_inc(g_a, i, j, inc)
C            long NGA_Read_inc(int g_a, int subscript[], long inc)
C++          long GA::GlobalArray::readInc(int subscript[], long inc)

```

This function applies to integer arrays only. It atomically reads and increments an element in an integer array. It performs

a(subscripts) += inc

and returns the original value (before the update) of *a(subscript)*.

4.3 Scatter/Gather

Scatter and *gather* transfer a specified set of elements to and from global arrays. They are one-sided: that is they don't need the cooperation of the process(es) who own the referenced elements in the global array.

Scatter puts array elements into a global array, which is

```

n-D Fortran  subroutine nga_scatter(g_a, v, subsarray, n)
2-D Fortran  subroutine ga_scatter(g_a, v, i, j, n)
C            void NGA_Scatter(int g_a, void *v, int *subsarray[], int
                           n)
C++          void GA::GlobalArray::scatter(void *v, int *subsarray[],
                           int n)

```

It performs (in C notation)

```

for(k=0; k<= n; k++) {
    a[subsArray[k][0]][subsArray[k][1]][subsArray[k][2]]... = v[k];
}


```

Example:

Scatter the 5 elements into a 10x10 global array

Element 1	v[0] = 5	subsArray[0][0] = 2
		subsArray[0][1] = 3
Element 2	v[1] = 3	subsArray[1][0] = 3
		subsArray[1][1] = 4
Element 3	v[2] = 8	subsArray[2][0] = 8
		subsArray[2][1] = 5
Element 4	v[3] = 7	subsArray[3][0] = 3
		subsArray[3][1] = 7
Element 5	v[4] = 2	subsArray[4][0] = 6
		subsArray[4][1] = 3

After the scatter operation, the five elements would be scattered into the global array as shown in the following figure.



Gather is the reverse operation of *scatter*. It gets the array elements from a global array into a local array.

```

n-D Fortran  subroutine nga_gather(g_a, v, subsarray, n)
2-D Fortran  subroutine ga_gather(g_a, v, i, j, n)
C            void NGA_Gather(int g_a, void *v, int *subsarray[], int n)
C++          void GA::GlobalArray::gather(void *v, int *subsarray[], int
                                           n)
    
```

It performs (in C notation)

```

for(k=0; k<= n; k++){
    v[k] = a[subsArray[k][0]][subsArray[k][1]][subsArray[k][2]]...;
}
    
```

4.4 Periodic Interfaces


Periodic interfaces to the one-sided operations have been added to Global Arrays in **version 3.1** to support some computational fluid dynamics problems on multidimensional grids. They provide an index translation layer that allows to use put, get, and accumulate operations possibly extending beyond the boundaries of a global array. The references that are outside of the boundaries are wrapped up inside the global array. To better illustrate these operations, look the following example:

Example:

Assume a two dimensional global array `g_a` with dimensions 5 X 5.

	1	2	3	4	5
1	1	6	11	16	21
2	2	7	12	17	22
3	3	8	13	18	23
4	4	9	14	19	24
5	5	10	15	20	25

To access a patch $[2:4, -1:3]$, one can assume that the array is wrapped over in the second dimension, as shown in the following figure



Therefore the patch $[2:4, -1:3]$ is

```

17  22  2  7  12
18  23  3  8  13
19  24  4  9  14

```

Periodic operations extend the boudary of each dimension in two directions, toward lower bound and toward the upper bound. For any dimension with $lo(i)$ to $hi(i)$, where $1 < i < ndim$, it extends the range from

$[lo(i) : hi(i)]$


to

$[(lo(i)-1-(hi(i)-lo(i)+1)) : (lo(i)-1)], [lo(i) : hi(i)],$ and $[(hi(i)+1) : (hi(i)+1+(hi(i)-lo(i)+1))],$ or $[(lo(i)-1-(hi(i)-lo(i)+1)) : (hi(i)+1+(hi(i)-lo(i)+1))].$

Even though the patch span in a much large range, the length must always be less, or equals to $(hi(i)-lo(i)+1)$.

Example:

For a 2 x 2 array as shown in the following figure, where the dimensions are $[1:2, 1:2]$, periodic operations would look the range of each dimensions as $[-1:4, -1:4]$.



Current version of GA supports three periodic operations. They are

- periodic get,
- periodic put, and
- periodic acc.

Periodic Get copies data from a global array section to a local array, which is almost the same as regular *get*, except the indices of the patch can be outside the boundaries of each dimension.

```

Fortran  subroutine nga_periodic_get(g_a, lo, hi, buf, ld)
C        void NGA_Periodic_get(int g_a, int lo[], int hi[], void *buf,
                                int ld[])
C++      void GA::GlobalArray::periodicGet(int lo[], int hi[], void
                                *buf, int ld[])
  
```

Similar to regular *get*, *lo* and *hi* specify where the data should come from in the global array, and *ld* specifies the stride information of the local array *buf*.

Example:

Let us look at the first example in this section. It is 5 x 5 two dimensional global array. Assume that the local buffer is an 4x3 array.

	1	2	3	4	5
1	1	6	11	16	21
2	2	7	12	17	22
3	3	8	13	18	23
4	4	9	14	19	24
5	5	10	15	20	25

Also assume that

```

lo[0] = -1, hi[0] = 2,
lo[1] = 4,  hi[1] = 6,  and
ld[0] = 4
  
```

After the periodic get, the local buffer `buf` would be

```
19 24 4
20 25 5
16 21 1
17 22 2
```

Periodic Put is the reverse operations of *Periodic Get*. It copies data from the local array to the global array section, which is

```
Fortran  subroutine nga_periodic_put(g_a, lo, hi, buf, ld)
C        void NGA_Periodic_put(int g_a, int lo[], int hi[], void *buf,
                        int ld[])
C++      void GA::GlobalArray::periodicPut(int lo[], int hi[], void
                        *buf, int ld[])
```

Similar to regular *put*, `lo` and `hi` specify where the data should go in the global array; `ld` specifies the stride information of the local array `buf`.

Periodic Put/Get (also include the *Accumulate*, which will be discussed later in this section) divide the patch into several smaller patches. For those smaller patches that are outside the global array, adjust the indices so that they rotate back to the original array. After that call the regular *Put/Get/Accumulate*, for each patch, to complete the operations.

Example:

Look at the example for periodic get. Because it is a 5 x 5 global array, the valid indices for each dimension are

```
dimension 0: [1 : 5]
dimension 1: [1 : 5]
```

The specified `lo` and `hi` are apparently out of the range of each dimension:


```
dimension 0: [-1 : 2] --> [-1 : 0]  -- wrap back --> [4 : 5]
                        [ 1 : 2]  ok

dimension 1: [ 4 : 6] --> [ 4 : 5]  ok
                        [ 6 : 6]  -- wrap back --> [1 : 1]
```

Hence, there will be four smaller patches after the adjustment. They are

```
patch 0:  [4 : 5, 4 : 5]
patch 1:  [4 : 5, 1 : 1]
patch 2:  [1 : 2, 4 : 5]
patch 3:  [1 : 2, 1 : 1]
```

as shown in the following figure



Of course the destination addresses of each smaller patch in the local buffer also need to be calculated.

Similar to regular *Accumulate*, *Periodic Accumulate* combines the data from the local array with data in the global array section, which is

```

Fortran  subroutine nga_periodic acc(g_a, lo, hi, buf, ld, alpha)
C        void NGA_Periodic acc(int g_a, int lo[], int hi[], void *buf,
                                int ld[],
                                void *alpha)
C++      void GA::GlobalArray::periodicAcc(int lo[], int hi[], void
                                *buf, int ld[],
                                void *alpha)

```

The local array is assumed to have the same number of dimensions as the global array. Users don't need to worry about where the region defined by `lo` and `hi` is physically located. The function performs

*global array section (lo[], hi[]) += alpha * buf*

Example:

Let us look at the same example as above. There is 5 x 5 two dimensional global array. Assume that the local buffer is an 4x3 array.

	1	2	3	4	5
1	1	6	11	16	21
2	2	7	12	17	22
3	3	8	13	18	23
4	4	9	14	19	24
5	5	10	15	20	25

Also assume that

```

lo[0] = -1, hi[0] = 2,
lo[1] = 4,  hi[1] = 6,  and
ld[0] = 4.

```

The local buffer `buf` is

```

1  5  9
4  6  5
3  2  1
7  8  2

```

and the $\alpha = 2$.

After the *Periodic Accumulate* operation, the global array will be

	1	2	3	4	5
1	3	6	11	22	25
2	6	7	12	24	38
3	3	8	13	18	23
4	22	9	14	21	34
5	15	10	15	28	37

5. Interprocess Synchronization

Global Arrays provide three types of synchronization calls to support different synchronization styles.

- Lock with mutex:* is useful for a shared memory model. One can lock a mutex, to exclusively access a critical section.
- Fence:* guarantees that the Global Array operations issued from the calling process are complete. The fence operation is local.
- Sync:* is a barrier. It synchronizes processes and ensures that all Global Array operations completed. Sync operation is collective.

5.1 Lock and Mutex

Lock works together with mutex. It is a simple synchronization mechanism used to protect a critical section. To enter a critical section, typically, one needs to do:

1. *Create mutexes*
2. *Lock on a mutex*
3. ...
Do the exclusive operation in the critical section
- ...
4. *Unlock the mutex*
5. *Destroy mutexes*

The function

```
Fortran  logical function ga_create_mutexes(number)
C        int GA_Create_mutexes(int number)
C++      int GA::GAServices::createMutexes(int number)
```

creates a set containing the *number* of mutexes. Only one set of mutexes can exist at a time. Mutexes can be created and destroyed as many times as needed. Mutexes are numbered: 0, ..., *number*-1.

The function

```
Fortran  logical function ga_destroy_mutexes()
C        int GA_Destroy_mutexes()
C++      int GA::GAServices::destroyMutexes()
```

destroys the set of mutexes created with `ga_create_mutexes`.

Both `ga_create_mutexes` and `ga_destroy_mutexes` are collective operations.

The functions

```
Fortran  subroutine ga_lock(int mutex)
          subroutine ga_unlock(int mutex)
C        void GA_lock(int mutex)
          void GA_unlock(int mutex)
```



```
C++      void GA::GAServices::lock(int mutex)
          void GA::GAServices::unlock(int mutex)
```

lock and unlock a mutex object identified by the `mutex` number, respectively. It is a fatal error for a process to attempt to lock a mutex which has already been locked by this process, or unlock a mutex which has not been locked by this process.

Example 1:

Use one mutex and the lock mechanism to enter the critical section.

```
status = ga_create_mutexes(1)
if(.not.status) then
    call ga_error('ga_create_mutexes failed ',0)
endif
call ga_lock(0)

... do something in the critical section
call ga_put(g_a, ...)
...

call ga_unlock(0)
if(.not.ga_destroy_mutexes()) then
    call ga_error('mutex not destroyed',0)
```

5.2 Fence

Fence blocks the calling process until all the data transfers corresponding to the Global Array operations initiated by this process complete. The typical scenario that it is being used is

1. *Initialize the fence*
2. ...
 Global array operations
- ...
3. *Fence*

This would guarantee the operations between step 1 and 3 are complete.

The function

```
Fortran  subroutine ga_init_fence()
C        void GA_Init_fence()
C++      void GA::GAServices::initFence()
```

Initializes tracing of completion status of data movement operations.

The function

```
Fortran  subroutine ga_fence()
```

```

C          void GA_Fence()
C++       void GA::GAServices::fence()

```

blocks the calling process until all the data transfers corresponding to GA operations called after `ga_init_fence` complete.

`ga_fence` must be called after `ga_init_fence`. A barrier, `ga_sync`, assures completion of all data transfers and implicitly cancels outstanding `ga_init_fence`. `ga_init_fence` and `ga_fence` must be used in pairs, multiple calls to `ga_fence` require the same number of corresponding `ga_init_fence` calls. `ga_init_fence/ga_fence` pairs can be nested.

Example 1:

Since `ga_put` might return before the data reaches the final destination `ga_init_fence` and `ga_fence` allow the process to wait until the data is actually moved:

```

call ga_init_fence()
call ga_put(g_a, ...)
call ga_fence()

```

Example 2:

`ga_fence` works for multiple GA operations.

```

call ga_init_fence()
call ga_put(g_a, ...)
call ga_scatter(g_a, ...)
call ga_put(g_b, ...)
call ga_fence()

```

The calling process will be blocked until data movements initiated by two calls to `ga_put` and one `ga_scatter` complete.

5.3 Sync

Sync is a collective operation. It acts as a barrier, which synchronizes all the processes and ensures that all the Global Array operations are complete at the call.

The function is

```

Fortran  subroutine ga_sync()
C        void GA_Sync()
C++     void GA::GAServices::sync()

```

Sync should be inserted as necessary. With many sync calls, the application performance would suffer.

6. Collective Array Operations

Global Arrays provide functions for collective array operations, targeting both whole arrays and patches (portions of global arrays). Collective operations require all the processes to make the call. In the underlying implementation, each process deals with its local data. These functions include:

- basic array operations,
- linear algebra operations, and
- interfaces to third party software packages.

6.1 Basic Array Operations

Global Arrays provide several mechanisms to manipulate contents of the arrays. One can set all the elements in an array/patch to a specific value, or as a special case set to zero. Since GA does not explicitly initialize newly created arrays, these calls are useful for initialization of an array/patch. (To fill the array with different values for each element, one can choose the one sided operation *putor* each process can initialize its local portion of an array/patch like ordinary local memory). One can also scale the array/patch by a certain factor, or copy the contents of one array/patch to another.

6.1.1 Whole Arrays

These functions apply to the entire array.

The function

```
Fortran  subroutine ga_zero(g_a)
C        void GA_Zero(int g_a)
C++     void GA::GlobalArray::zero()
```

sets all the elements in the array to zero.

To assign a single value to all the elements in an array, use the function

```
Fortran  subroutine ga_fill(g_a, val)
C        void GA_Fill(int g_a, void *val)
C++     void GA::GlobalArray::fill(void *val)
```

It sets all the elements in the array to the value *val*. The *val* must have the same data type as that of the array.

The function

```
Fortran  subroutine ga_scale(g_a, val)
C        void GA_Scale(int g_a, void *val)
C++     void GA::GlobalArray::scale(void *val)
```

scales all the elements in the array by factor *val*. Again the *val* must be the same data type as that of the array itself.

The above three functions are dealing with one global array, to set values or change all the elements together. The following functions are for copying data between two arrays.

The function

```

Fortran subroutine ga_copy(g_a, g_b)
C         void GA_Copy(int g_a, int g_b)
C++       void GA::GlobalArray::copy(const GA::GlobalArray * g_a)

```

copies the contents of one array to another. The arrays must be of the same data type and have the same number of elements.

For global arrays containing ghost cells, the ghost cell data can be filled in with the corresponding data from neighboring processors using the command

```

n-d Fortran subroutine ga_copy(g_a, g_b)
C         void GA_Copy(int g_a, int g_b)
C++       void GA::GlobalArray::copy(const GA::GlobalArray * g_a)

```

```

n-d Fortran subroutine ga_update_ghosts(g_a)
C         void GA_Update_ghosts(int g_a)
C++       void GA::GlobalArray::updateGhosts()

```

This operation updates the ghost cell data by assuming periodic, or wrap-around, boundary conditions similar to those described for the `nga_periodic_get` operations described above. The wrap-around conditions are always applied, it is up to the individual application to decide whether or not the data in the ghost cells should be used. The update operation is illustrated below for a simple 4x2 global array distributed across two processors. The ghost cells are one element wide in each dimension.

Global Array with Uninitialized Ghost Cells

*	*	*	*
*	1	2	*
*	5	6	*
*	*	*	*

*	*	*	*
*	3	4	*
*	7	8	*
*	*	*	*

Global Array after Updating Ghost Cells

8	5	6	7
4	1	2	3
8	5	6	7
4	1	2	3

6	7	8	5
2	3	4	1
6	7	8	5
2	3	4	1

```

n-d Fortran logical function nga_update_ghosts_dir(g_a, dimension,
idir, flag)

```

```

C          int NGA_Update_ghosts_dir(int g_a, int dimension, int
idir, int cflag)
C++       int GA::GlobalArray::updateGhostsDir(int dimension, int
idir, int cflag)

```

This function can be used to update the ghost cells along individual directions.

It is designed for algorithms that can overlap updates with computation. The variable dimension indicates which coordinate direction is to be updated (e.g. dimension = 1 would correspond to the y axis in a two or three dimensional system), the variable idir can take the values ± 1 and indicates whether the side that is to be updated lies in the positive or negative direction, and cflag indicates whether or not the corners on the side being updated are to be included in the update. The following calls would be equivalent to a call to `GA_Update_ghosts` for a 2-dimensional system:

```

status = NGA_Update_ghost_dir(g_a,0,-1,1);
status = NGA_Update_ghost_dir(g_a,0,1,1);
status = NGA_Update_ghost_dir(g_a,1,-1,0);
status = NGA_Update_ghost_dir(g_a,1,1,0);

```

The variable cflag is set equal to 1 (or non-zero) in the first two calls so that the corner ghost cells are update, it is set equal to 0 in the second two calls to avoid redundant updates of the corners. Note that updating the ghosts cells using several independent calls to the `nga_update_ghost_dir` functions is generally not as efficient as using `GA_Update_ghosts` unless the individual calls can be effectively overlapped with computation. This is a collective operation.

6.1.2 Patches

GA provides a set of operations on segments of the global arrays, namely patch operations. These functions are more general, in a sense they can apply to the entire array(s). As a matter of fact, many of the Global Array collective operations are based on the patch operations, for instance, the `GA_Print` is only a special case of `NGA_Print_patch`, called by setting the bounds of the patch to the entire global array. There are two interfaces for Fortran, one for two dimensional and the other for n-dimensional (one to seven). The (n-dimensional) interface can surely handle the two dimensional case as well. It is available for backward compatibility purposes. The functions dealing with n-dimensional patches use the "nga" prefix and those dealing with two dimensional patches start with the "ga" prefix.

The function

```

Fortran    subroutine nga_zero_patch(g_a, alo, ahi)
C          void NGA_Zero_patch(int g_a, int lo[] int hi[])
C++       void GA::GlobalArray::zeroPatch(int lo[] int hi[])

```

is similar to `ga_zero`, except that instead of applying to entire array, it sets only the region defined by *lo* and *hi* to zero.

One can assign a single value to all the elements in a patch with the function:

```

n-DFortran subroutine nga_fill_patch(g_a, lo, hi, val)
2-DFortran subroutine ga_fill_patch(g_a, ilo, ihi, jlo, jhi, val)
C          void NGA_Fill_patch(int g_a, int lo[] int hi[], void *val)
C++       void GA::GlobalArray::fillPatch(int lo[] int hi[], void

```

*val)

The lo and hi defines the patch and the val is the value to set.

The function

```

n-DFortran  subroutine nga_scale_patch(g_a, lo, hi, val)
2-DFortran  subroutine ga_scale_patch(g_a, ilo, ihi, jlo, jhi, val)
C           void NGA_Scale_patch(int g_a, int lo[], int hi[], void *val)
C++        void GA::GlobalArray::scalePatch(int lo[], int hi[], void
                                           *val)

```

scales the patch defined by lo and hi by the factor val.

The copy patch operation is one of the fundamental and frequently used functions. The function

```

n-DFortran  subroutine nga_copy_patch(trans, g_a, alo, ahi,
                                           g_b, blo, bhi)
2-DFortran  subroutine ga_copy_patch(trans, g_a, ailo, aihi, ajlo,
                                           ajhi, g_b, bilo, bihi, bjlo, bjhi)
C           void NGA_Copy_patch(char trans, int g_a, int alo[], int
                                           ahi[],
                                           int g_b, int blo[], int bhi[])
C++        void GA::GlobalArray::copyPatch(char trans, const
                                           GA::GlobalArray* g_a,
                                           int alo[], int ahi[], int blo[], int bhi[])

```

copies one patch defined by alo and ahi in one global array g_a to another patch defined by blo and bhi in another global array g_b. The current implementation requires that the source patch and destination patch must be on different global arrays. They must also be the same data type. The patches may be of different shapes, but the number of elements must be the same. During the process of copying, the transpose operation can be performed by specifying trans.

Example: Assume that there two 8x6 Global Arrays, g_a and g_b, distributed on three processes. The operation of nag_copy_patch(Fortran notation), from

```
g_a: alo = {2, 2}, ahi = {4, 5}
```


to

```
g_b: blo = {3, 4}, bhi = {6, 6}
```


and

```
trans = 0
```

involves reshaping. It is illustrated in the following figure.



One step further, if one also want to perform the transpose operation during the copying, *i.e.* set `trans = 1`, it will look like:



If there is no reshaping or transpose, the operation can be fast (internally calling `nga_put`). Otherwise, it would be slow (internally calling `nga_scatter`, where extra time is spent on preparing the indices). Also note that extra memory is required to hold the indices if the operation involves reshaping or transpose.

6.2 Linear Algebra

Global arrays provide three linear algebra operations: addition, multiplication, and dot product. There are two sets of functions, one for the whole array and the other for the patches.

6.2.1 Whole Arrays

The function

```
Fortran subroutine ga_add(alpha, g_a, beta, g_b, g_c)
C      void GA_Add(void *alpha, int g_a, void *beta, int g_b, int g_c)
C++    void GA::GlobalArray::add(void *alpha, const GA::GlobalArray*
                                   g_a,
                                   void *beta, const GA::GlobalArray* g_b)
```

adds two arrays, `g_a` and `g_b`, and saves the results to `g_c`. The two source arrays can be scaled by certain

factors. This operation requires the two source arrays have the same number of elements and the same data types, but the arrays can have different shapes or distributions. `g_ccan` also be `g_a` or `g_b`. It is encouraged to use this function when the two source arrays are identical in distributions and shapes, because of its efficiency. It would be less efficient if the two source arrays are different in distributions or shapes.

Matrix multiplication operates on two matrices, therefore the array must be two dimensional. The function

```

Fortran subroutine ga_dgemm(transa, transb, m, n, k,
                                alpha, g_a, g_b, beta, g_c )
C      void GA_Dgemm(char ta, char tb, int m, int n, int k,
                        double alpha, int g_a, int g_b,
                        double beta, int g_c )
C++    void GA::GlobalArray::dgemm(char ta, char tb, int m, int n, int
                                    k,
                                    double alpha, const GA::GlobalArray* g_a,
                                    const GA::GlobalArray* g_b, double beta)

```

Performs one of the matrix–matrix operations:

$$C := \alpha * op(A) * op(B) + \beta * C,$$

where $op(X)$ is one of

$$op(X) = X \text{ or } op(X) = X',$$

α and β are scalars, and A , B and C are matrices, with $op(A)$ an m by k matrix, $op(B)$ a k by n matrix and C an m by n matrix.

On entry, `transa` specifies the form of $op(A)$ to be used in the matrix multiplication as follows:

`ta = 'N' or 'n', $op(A) = A$.`

`ta = 'T' or 't', $op(A) = A'$.`

The function

```

Fortran integer      function ga_idot(g_a, g_b)
double precision function ga_ddot(g_a, g_b)
double complex  function ga_zdot(g_a, g_b)
C      long      GA_Idot(int g_a, int g_b)
double      GA_Ddot(int g_a, int g_b)
DoubleComplex GA_Zdot(int g_a, int g_b)
C++    long      GA::GlobalArray::idot(const GA::GlobalArray*
                                          g_a)
double      GA::GlobalArray::ddot(const GA::GlobalArray*
                                          g_a)
DoubleComplex GA::GlobalArray::zdot(const GA::GlobalArray*
                                          g_a)

```

computes the element–wise dot product of two arrays. It is available as three separate functions, corresponding to *integer*, *double precision* and *double complex* data types.

The following functions apply to the 2-dimensional whole arrays only. There are no corresponding functions for patch operations.

The function

```
Fortran  subroutine ga_symmetrize(g_a)
C        void GA_Symmetrize(int g_a)
C++      void GA::GlobalArray::symmetrize()
```

symmetrizes matrix A represented with handle g_a: $A = .5 * (A + A')$.

The function

```
Fortran  subroutine ga_transpose(g_a, g_b)
C        void GA_Transpose(int g_a, int g_b)
C++      void GA::GlobalArray::transpose(const GA::GlobalArray* g_a)
```

transposes a matrix: $B = A'$.

6.2.2 Patches

The functions

```
n-DFortran  subroutine nga_add_patch(alpha, g_a, alo, ahi,
                                   beta, g_b, blo, bhi,
                                   g_c, clo, chi)
2-DFortran  subroutine ga_add_patch(alpha, g_a, ailo, aihi, ajlo,
                                   ajhi,
                                   beta, g_b, bilo, bihi, bjlo,
                                   bjhi,
                                   g_c, cilo, cihi, cjlo,
                                   cjhi)
C           void NGA_Add_patch(void *alpha, int g_a, int alo[], int
                                   ahi[],
                                   void *beta, int g_b, int blo[], int
                                   bhi[],
                                   int g_c, int clo[], int
                                   chi[])
C++        void GA::GlobalArray::addPatch(void *alpha, const
                                   GA::GlobalArray* g_a,
                                   int alo[], int ahi[], void *beta, const
                                   GA::GlobalArray* g_b,
                                   int blo[], int bhi[], int clo[], int
                                   chi[])
```

add element-wise two patches and save the results into another patch. Even though it supports the addition of two patches with different distributions or different shapes (the number of elements must be the same), the operation can be expensive, because there can be extra copies which effect memory consumption. The two source patches can be scaled by a factor for the addition. The function is smart enough to detect the case that the patches are exactly the same but the global arrays are different in shapes. It handles the case as if for the

arrays were identically distributed, thus the performance will not suffer.

The matrix multiplication is the only operation on array patches that is restricted to the two dimensional domain, because of its nature. It works for *double* and *double complex* data types. The prototype is

```

Fortran subroutine ga_matmul_patch(transa, transb, alpha, beta,
                                     g_a, ailo, aihi, ajlo, ajhi,
                                     g_b, bilo, bihi, bjlo, bjhi,
                                     g_c, cilo, cihi, cjlo, cjhi)
C      void GA_Matmul_patch(char *transa, char* transb, void* alpha,
                             void *beta,
                             int g_a, int ailo, int aihi, int ajlo, int
                             ajhi,
                             int g_b, int bilo, int bihi, int bjlo, int
                             bjhi,
                             int g_c, int cilo, int cihi, int cjlo, int
                             cjhi)
C++    void GA::GlobalArray::matmulPatch(char *transa, char* transb,
                                             void* alpha, void *beta,
                                             const GlobalArray * g_a, int ailo, int aihi,
                                             int ajlo, int ajhi,
                                             const GlobalArray * g_b, int bilo, int bihi,
                                             int bjlo, int bjhi,
                                             int cilo, int cihi, int cjlo, int cjhi)

```

It performs

$$C[cilo:cihi,cjlo:cjhi] := \alpha * AA[ailo:aihi,ajlo:ajhi] * BB[bilo:bihi,bjlo:bjhi] + \beta * C[cilo:cihi,cjlo:cjhi]$$

where $AA = op(A)$, $BB = op(B)$, and $op(X)$ is one of

$$op(X) = X \text{ or } op(X) = X',$$

Valid values for transpose argument: 'n', 'N', 't', 'T'.

The dot operation computes the element-wise dot product of two (possibly transposed) patches. It is implemented as three separate functions, corresponding to *integer*, *double precision* and *double complex* data types. They are

```

n-DFortran integer function nga_idot_patch(g_a, ta, alo, ahi,
                                             g_b, tb, blo, bhi)
double precision function nga_ddot_patch(g_a, ta, alo, ahi,
                                             g_b, tb, blo, bhi)
double complex function nga_zdot_patch(g_a, ta, alo, ahi,
                                             g_b, tb, blo, bhi)

2-DFortran integer function ga_idot_patch(g_a, ta, ailo, aihi,
                                             ajlo, ailo, g_b, tb, bilo, bihi, bjlo, bjhi)
double precision function ga_ddot_patch(g_a, ta, ailo,
aihi,

```

Global Arrays User Manual

```
        ajlo, ailo, g_b, tb, bilo, bihi, bjlo, bjhi)
double complex function ga_zdot_patch(g_a, ta, ailo, aihi,
        ajlo, ailo, g_b, tb, bilo, bihi, bjlo, bjhi)
```

```
C      Integer NGA_Idot_patch(int g_a, char* ta, int alo[], int
        ahi[],
        int g_b, char* tb, int blo[], int bhi[])
double NGA_Ddot_patch(int g_a, char* ta, int alo[], int
        ahi[],
        int g_b, char* tb, int blo[], int bhi[])
DoubleComplex NGA_Zdot_patch(int g_a, char* ta, int alo[],
        int ahi[],
        int g_b, char* tb, int blo[], int bhi[])
```

```
C++    Integer GA::GlobalArray::idotPatch(const GA::GlobalArray*
        g_a,
        char* ta, int alo[], int ahi[],
        char* tb, int blo[], int bhi[])
double GA::GlobalArray::ddotPatch(const GA::GlobalArray*
        g_a,
        char* ta, int alo[], int ahi[],
        char* tb, int blo[], int bhi[])
DoubleComplex GA::GlobalArray::zdotPatch(const
        GA::GlobalArray* g_a,
        char* ta, int alo[], int ahi[],
        char* tb, int blo[], int bhi[])
```

The patches should be of the same data types and have the same number of elements. Like the array addition, if the source patches have different distributions/shapes, or it requires transpose, the operation would be less efficient, because there could be extra copies and/or memory consumption.

6.2.3 Element-wise operations

These operations work on individual array elements rather than arrays as matrices in the sense of linear algebra operations. For example multiplication of elements stored in arrays is a completely different operation than matrix multiplication.

Fortran subroutine ga_abs_value(g_a)

C void GA_Abs_value(int g_a)

C++ void GA::GlobalArray::absValue(int g_a)

Take element-wise absolute value of the array.

Fortran subroutine ga_abs_value_patch(g_a, lo, hi)

C void GA_Abs_value_patch(int g_a, int lo[], int hi[])

C++ void GA::GlobalArray::absValuePatch(int lo[], int hi[])

Take element-wise absolute value of the patch.

```

Fortran subroutine ga_add_constant(g_a, alpha)
C      void GA_Add_constant(int g_a, void* alpha)
C++    void GA::GlobalArray::addConstant(void* alpha)

```

Add the constant pointed by alpha to each element of the array.

```

Fortran subroutine ga_add_constant_patch(g_a, lo, hi, alpha)
C      void GA_Add_constant_patch(int g_a, int lo[], int hi[], void* alpha)
C++    void GA::GlobalArray::addConstantPatch(void* alpha)

```

Add the constant pointed by alpha to each element of the patch.

```

Fortran subroutine ga_recip(g_a)
C      void GA_Recip(int g_a)
C++    void GA::GlobalArray::recip()

```

Take element-wise reciprocal of the array.

```

Fortran subroutine ga_recip_patch(g_a, lo, hi)
C      void GA_Recip_patch(int g_a, int lo[], int hi[])
C++    void GA::GlobalArray::recipPatch(int lo[], int hi[])

```

Take element-wise reciprocal of the patch.

```

Fortran subroutine ga_elem_multiply(g_a, g_b, g_c)
C      void GA_Elem_multiply(int g_a, int g_b, int g_c)
C++    void GA::GlobalArray::elemMultiply(const GA::GlobalArray * g_a,
                                           const GA::GlobalArray * g_b)

```

Computes the element-wise product of the two arrays which must be of the same types and same number of elements. For two-dimensional arrays,

$$c(i, j) = a(i, j) * b(i, j)$$

The result (c) may replace one of the input arrays (a/b).

```

Fortran subroutine ga_elem_multiply_patch(g_a, alo, ahi, g_b, blo, bhi, g_c, clo, chi)
C      void GA_Elem_multiply_patch(int g_a, int alo[], int ahi[], int g_b, int blo[],
                                     int bhi[], int g_c, int clo[], int chi[])
C++    void GA::GlobalArray::elemMultiplyPatch(const GA::GlobalArray * g_a,
                                                  int alo[], int ahi[],
                                                  const GA::GlobalArray * g_b, int blo[],
                                                  int bhi[], int clo[], int chi[])

```

Computes the element-wise product of the two patches which must be of the same types and same number of elements. For two-dimensional arrays,

$$c(i, j) = a(i, j) * b(i, j)$$

The result (c) may replace one of the input arrays (a/b).

Fortran subroutine ga_elem_divide(g_a, g_b, g_c)
C void GA_Elem_divide(Integer g_a, Integer g_b, Integer g_c)
C++ void GA::GlobalArray::elemDivide(const GA::GlobalArray * g_a,
const GA::GlobalArray * g_b)

Computes the element-wise quotient of the two arrays which must be of the same types and same number of elements. For two-dimensional arrays,

$$c(i, j) = a(i, j) / b(i, j)$$

The result (c) may replace one of the input arrays (a/b). If one of the elements of array g_b is zero, the quotient for the element of g_c will be set to GA_NEGATIVE_INFINITY.

Fortran subroutine ga_elem_divide_patch(g_a, alo, ahi, g_b, blo, bhi, g_c, clo, chi)
C void GA_Elem_divide_patch(int g_a, int alo[], int ahi[], int g_b, int blo[],
int bhi[], int g_c, int clo[], int chi[])
C++ void GA::GlobalArray::elemDividePatch(const GA::GlobalArray * g_a,
int alo[], int ahi[],
const GA::GlobalArray * g_b, int blo[],
int bhi[], int clo[], int chi[])

Computes the element-wise quotient of the two patches which must be of the same types and same number of elements. For two-dimensional arrays,

$$c(i, j) = a(i, j) / b(i, j)$$

The result (c) may replace one of the input arrays (a/b).

Fortran subroutine ga_elem_maximum(g_a, g_b, g_c)
C void GA_Elem_maximum(Integer g_a, Integer g_b, Integer g_c)
C++ void GA::GlobalArray::elemMaximum(const GA::GlobalArray * g_a,
const GA::GlobalArray * g_b)

Computes the element-wise maximum of the two arrays which must be of the same types and same number of elements. For two dimensional arrays,

$$c(i, j) = \max\{a(i, j), b(i, j)\}$$

The result (c) may replace one of the input arrays (a/b).

Fortran subroutine ga_elem_maximum_patch(g_a, alo, ahi, g_b, blo, bhi, g_c, clo, chi)
C void GA_Elem_maximum_patch(int g_a, int alo[], int ahi[], int g_b, int blo[],
int bhi[], int g_c, int clo[], int chi[])
C++ void GA::GlobalArray::elemMaximumPatch(const GA::GlobalArray * g_a,
int alo[], int ahi[],
const GA::GlobalArray * g_b, int blo[],
int bhi[], int clo[], int chi[])

Computes the element-wise maximum of the two patches
which must be of the same types and same number of
elements. For two-dimensional of noncomplex arrays,

$$c(i, j) = \max\{a(i, j), b(i, j)\}$$

If the data type is complex, then

$$c(i, j).real = \max\{|a(i, j)|, |b(i, j)|\} \text{ while } c(i, j).image = 0.$$

The result (c) may replace one of the input arrays (a/b).

Fortran subroutine ga_elem_minimum(g_a, g_b, g_c)
C void GA_Elem_minimum(Integer g_a, Integer g_b, Integer g_c);
C++ void GA::GlobalArray::elemMinimum(const GA::GlobalArray * g_a,
const GA::GlobalArray * g_b)

Computes the element-wise minimum of the two arrays
which must be of the same types and same number of
elements. For two dimensional arrays,

$$c(i, j) = \min\{a(i, j), b(i, j)\}$$

The result (c) may replace one of the input arrays (a/b).

Fortran subroutine ga_elem_minimum_patch(g_a, alo, ahi, g_b, blo, bhi, g_c, clo, chi)
C void GA_Elem_minimum_patch(int g_a, int alo[], int ahi[], int g_b,
int blo[], int bhi[], int g_c, int clo[], int chi[])
C++ void GA::GlobalArray::elemMinimumPatch(const GA::GlobalArray * g_a,
int alo[], int ahi[],
const GA::GlobalArray * g_b, int blo[],
int bhi[], int clo[], int chi[])

Computes the element-wise minimum of the two patches
which must be of the same types and same number of
elements. For two-dimensional of noncomplex arrays,

$$c(i, j) = \min\{a(i, j), b(i, j)\}$$

If the data type is complex, then

$$c(i, j).real = \min\{|a(i, j)|, |b(i, j)|\} \text{ while } c(i, j).image = 0.$$

The result (c) may replace one of the input arrays (a/b).

```
Fortran subroutine ga_shift_diagonal(g_a, c)
C      void GA_Shift_diagonal(int g_a, void *c)
C++    void GA::GlobalArray::shiftDiagonal(void *c)
```

Adds this constant to the diagonal elements of the matrix.

```
Fortran subroutine ga_set_diagonal(g_a, g_v)
C      void GA_Set_diagonal(int g_a, int g_v)
C++    void GA::GlobalArray::setDiagonal(const GA::GlobalArray * g_v)
```

Sets the diagonal elements of this matrix g_a with the elements of the vector g_v.

```
Fortran subroutine ga_zero_diagonal( g_a)
C      void GA_Zero_diagonal(int g_a)
C++    void GA::GlobalArray::zeroDiagonal()
```

Sets the diagonal elements of this matrix g_a with zeros.

```
Fortran subroutine ga_add_diagonal(g_a, g_v)
C      void GA_Add_diagonal(int g_a, int g_v)
C++    void GA::GlobalArray::addDiagonal(const GA::GlobalArray * g_v)
```

Adds the elements of the vector g_v to the diagonal of this matrix g_a.

```
Fortran subroutine ga_get_diag(g_a, g_v)
C      void GA_Get_diag(int g_a, int g_v)
C++    void GA::GlobalArray::getDiagonal(const GA::GlobalArray * g_v)
```

Inserts the diagonal elements of this matrix g_a into the vector g_v.

```
Fortran subroutine ga_scale_rows( g_a, g_v)
C      void GA_Scale_rows(int g_a, int g_v)
C++    void GA::GlobalArray::scaleRows(const GA::GlobalArray * g_v)
```

Scales the rows of this matrix g_a using the vector g_v.

```
Fortran subroutine ga_scale_cols(g_a, g_v)
C      void GA_Scale_cols(int g_a, int g_v)
C++    void GA::GlobalArray::scaleCols(const GA::GlobalArray * g_v)
```

Scales the columns of this matrix g_a using the vector g_v.

Fortran subroutine ga_norm1(g_a, nm)
C void GA_Norm1(int g_a, double *nm)
C++ void GA::GlobalArray::norm1(double *nm)

Computes the 1–norm of the matrix or vector g_a.

Fortran subroutine ga_norm_infinity(g_a, nm)
C void GA_Norm_infinity(int g_a, double *nm)
C++ void GA::GlobalArray::normInfinity(double *nm)

Computes the 1–norm of the matrix or vector g_a.

Fortran subroutine ga_median(g_a, g_b, g_c, g_m)
C void GA_Median(int g_a, int g_b, int g_c, int g_m)
C++ void GA::GlobalArray::median(const GA::GlobalArray * g_a,
const GA::GlobalArray * g_b,
const GA::GlobalArray * g_c)

Computes the componentwise Median of three arrays g_a, g_b, and g_c, and stores the result in this array g_m. The result (m) may replace one of the input arrays (a/b/c).

Fortran subroutine ga_median_patch(g_a, alo, ahi, g_b, blo, bhi, g_c,
clo, chi, g_m, mlo, mhi)
C void GA_Median_patch(int g_a, int alo[], int ahi[], int g_b, int blo[],
int bhi[], int g_c, int clo[], int chi[], int g_m,
int mlo[], int mhi[])
C++ void GA::GlobalArray::medianPatch(const GA::GlobalArray * g_a, int alo[], int ahi[],
const GA::GlobalArray * g_b, int blo[], int bhi[],
const GA::GlobalArray * g_c, int clo[], int chi[],
int mlo[], int mhi[])

Computes the componentwise Median of three patches g_a, g_b, and g_c, and stores the result in this patch g_m. The result (m) may replace one of the input patches (a/b/c).

Fortran subroutine ga_step_max(g_a, g_b, step)
C void GA_Step_max(int g_a, int g_b, double *step)
C++ void GA::GlobalArray::stepMax(const GA::GlobalArray *g_a, double *step)

Calculates the largest multiple of a vector g_b that can be added to this vector g_a while keeping each element of this vector nonnegative.

Fortran subroutine ga_step_max2(g_xx, g_vv, g_xxll, g_xxuu, step2)
C void GA_Step_max2(int g_xx, int g_vv, int g_xxll, int g_xxuu, double *step2)
C++ void GA::GlobalArray::stepMax2(const GA::GlobalArray *g_vv,
const GA::GlobalArray *g_xxll,


```
const GA::GlobalArray *g_xxuu, double *step2)
```

Calculates the largest step size that should be used in a projected bound line search.

```
Fortran subroutine ga_step_max_patch(g_a, alo, ahi, g_b, blo, bhi, step)
C      void GA_Step_max_patch(int g_a, int *alo, int *ahi, int g_b, int *blo,
                                int *bhi, double *step)
C++    void GA::GlobalArray::stepMaxPatch(int *alo, int *ahi, const GA::GlobalArray * g_b,
                                            int *blo, int *bhi, double *step)
```

Calculates the largest multiple of a vector `g_b` that can be added to this vector `g_a` while keeping each element of this vector nonnegative.

```
Fortran subroutine ga_step_max2_patch( g_xx, xxlo, xxhi, g_vv, vvlo, vvhi, g_xll,
                                         xxlllo, xllhi, g_xxuu, xxuulo, xxuuhi, step2)
C      void GA_Step_max2_patch(int g_xx, int *xxlo, int *xxhi, int g_vv,
                                int *vvlo, int *vvhi, int g_xll, int *xlllo, int *xllhi,
                                int g_xxuu, int *xxuulo, int *xxuuhi, double *step2)
C++    void GA::GlobalArray::stepMax2Patch(int *xxlo, int *xxhi,
                                             const GA::GlobalArray * g_vv, int *vvlo, int *vvhi,
                                             const GA::GlobalArray * g_xll, int *xlllo, int *xllhi,
                                             const GA::GlobalArray * g_xxuu, int *xxuulo,
                                             int *xxuuhi, double *step2)
```

Calculates the largest step size that should be used in a projected bound line search.

6.3 Interfaces to Third Party Software Packages

There are many existing software packages designed for solving engineering problems. They are specialized in one or two problem domains, such as solving linear systems, eigen-vectors, and differential equations, etc. Global Arrays provide interfaces to several of these packages.

6.3.1 Scalapack

Scalapack is a well known software library for linear algebra computations on distributed memory computers. Global Arrays uses this library to solve systems of linear equations and also to invert matrices.

The function

```
Fortran integer function ga_solve(g_a, g_b)
C      int GA_Solve(int g_a, int g_b)
C++    int GA::GlobalArray::solve(const GA::GlobalArray * g_a)
```

solves a system of linear equations $A * X = B$. It first will call the Cholesky factorization routine and, if successful, will solve the system with the Cholesky solver. If Cholesky is not able to factorize A , then it will call the LU factorization routine and will solve the system with forward/backward substitution. On exit B will contain the solution X .

The function

```

Fortran integer function ga_llt_solve(g_a, g_b)
C       int GA_Llt_solve(int g_a, int g_b)
C++    int GA::GlobalArray::lltSolve(const GA::GlobalArray * g_a)

```

also solves a system of linear equations $A * X = B$, using the Cholesky factorization of an $N \times N$ double precision symmetric positive definite matrix A (handle g_a). On successful exit B will contain the solution X .

The function

```

Fortran subroutine ga_lu_solve(trans, g_a, g_b)
C       void GA_Lu_solve(char trans, int g_a, int g_b)
C++    void GA::GlobalArray::luSolve(char trans, const
GA::GlobalArray * g_a)

```

solves the system of linear equations $op(A)X = B$ based on the LU factorization. $op(A) = A$ or A' depending on the parameter `trans`. Matrix A is a general real matrix. Matrix B contains possibly multiple *rhs* vectors. The array associated with the handle g_b is overwritten by the solution matrix X .

The function

```

Fortran integer function ga_spd_invert(g_a)
C       int GA_Spd_invert(int g_a)
C++    int GA::GlobalArray::spdInvert( )

```

computes the inverse of a double precision matrix using the Cholesky factorization of a $N \times N$ double precision symmetric positive definite matrix A stored in the global array represented by g_a . On successful exit, A will contain the inverse.

6.3.2 PeIGS

The PeIGS library contains subroutines for solving standard and generalized real symmetric eigensystems. All eigenvalues and eigenvectors can be computed. The library is implemented using a message-passing model and is portable across many platforms. For more information and availability send a message to gi_fann@pnl.gov. Global Arrays use this library to solve eigen-value problems.

The function

```

Fortran subroutine ga_diag(g_a, g_s, g_v, eval)
C       void GA_Diag(int g_a, int g_s, int g_v, void *eval)
C++    void GA::GlobalArray::diag(const GA::GlobalArray* g_s,
                                const GA::GlobalArray* g_v, void *eval)

```

solves the generalized eigen-value problem returning all eigen-vectors and values in ascending order. The input matrices are not overwritten or destroyed.

The function

```

Fortran subroutine ga_diag_reuse(control, g_a, g_s, g_v, eval)
C       void GA_Diag_reuse(int control, int g_a, int g_s, int

```

```

g_v, void *eval)
    C++      void GA::GlobalArray::diagReuse(int control, const
GA::GlobalArray* g_s,
                                                    const GA::GlobalArray*g_v, void
*eval)

```

solves the generalized eigen-value problem returning all eigen-vectors and values in ascending order. Recommended for REPEATED calls if g_sis unchanged.

The function

```

Fortran  subroutine ga_diag_std(g_a, g_v, eval)
C        void GA_Diag_std(int g_a, int g_v, void *eval)
C++      void GA::GlobalArray::diagStd( const GA::GlobalArray* g_v,
void *eval)

```

solves the standard (non-generalized) eigenvalue problem returning all eigenvectors and values in the ascending order. The input matrix is neither overwritten nor destroyed.

6.3.3 Interoperability with Others

Global Arrays are interoperable with several other libraries, but do not provide direct interfaces for them. For example, one can make calls to and link with these libraries:

PETSc(the Portable, Extensible Toolkit for Scientific Computation) is developed by the Argonne National Laboratory. PETSc is a suite of data structures and routines for the scalable (parallel) solution of scientific applications modeled by partial differential equations. It employs the MPI standard for all message-passing communication, and is written in a data-structure-neutral manner to enable easy reuse and flexibility. Here is the instructions for using PETSc with GA.

CUMULVS (Collaborative User Migration User Library for Visualization and Steering) is developed by the Oak Ridge National Laboratory. CUMULVS is a software framework that enables programmers to incorporate fault-tolerance, interactive visualization and computational steering into existing parallel programs. Here is the instructions for using CUMULVS with GA.

6.4 Synchronization Control in Collective Operations

GA collective array operations are implemented by exploiting locality information to minimize or even completely avoid interprocessor communication or data copying. Before each processor accesses its own portion of the GA data we must assure that the data is in a consistent state. That means that there are no outstanding communication operations targeting that given global array portion pending while the data owner is accessing it. To accomplish that the GA collective array operations have implicit synchronization points: at the beginning and at the end of the operation. However, in many cases when collective array operations are called back-to-back or if the user does an explicit sync just before a collective array operation, some of the internal synchronization points could be merged or even removed if user can guarantee that the global array data is in the consistent state. The library offers a call for the user to eliminate the redundant synchronization points based on his/her knowledge of the application.

The function

```
Fortran  subroutine ga_mask_sync(prior_sync_mask,post_sync_mask)
C       void GA_Mask_sync(int prior_sync_mask,int post_sync_mask)
C++     void GA::GlobalArray::maskSync(int prior_sync_mask, int
post_sync_mask)
```

This operation should be used with a lot of care and only when the application code has been debugged and the user wishes to tune its performance. Making a call to this function with `prior_sync_mask` parameter set to false disables the synchronization done at the beginning of first collective array operation called after a call to this function. Similarly, making a call to this function by setting the `post_sync_mask` parameter to false disables the synchronization done at the ending of the first collective array operation called after a call to this function.

7. Utility Operations

Global Arrays include some utility functions to provide process, data locality, information, check the memory availability, etc. There are also several handy functions that print array distribution information, or summarize array usage information.

7.1 Locality Information

For a given global array element, or a given patch, sometimes it is necessary to find out who owns this element or patch. The function

```

n-DFortran logical functionnga_locate(g_a, subscript, owner)
2-DFortran logical functionga_locate(g_a, i, j, owner)
C          int_NGA_Locate(int g_a, int subscript[])
C++       int_GA::GlobalArray::locate(int subscript[])

```

tells who (process id) owns the elements defined by the array subscripts.

The function

```

n-DFortran logical functionnga_locate_region(g_a, lo, hi,
                                             map,proclist, np)
2-DFortran logical functionga_locate_region(g_a, ilo, ihi, jlo,jhi,
                                             map, np)
C          int_NGA_Locate_region(int g_a, int lo[], int hi[],int
                                *map[], int procs[])
C++       int_GA::GlobalArray::locateRegion(int lo[], int hi[],int
                                *map[], int procs[])

```

returns a list of GA process IDs that 'own' the patch.

The Global Arrays support an abstraction of a distributed array object. This object is represented by an integer handle. A process can access its portion of the data in the global array. To do this, the following steps need to be taken:

1. find the distribution of an array, which part of the data the calling process own
2. access the data
3. operate on the data: read/write
4. release the access to the data

The function

```

n-DFortran subroutinenga_distribution(g_a, iproc, lo, hi)
2-DFortran subroutinega_distribution(g_a, iproc, ilo, ihi, jlo, jhi)
C          void_NGA_Distribution(int g_a, int iproc, int lo[], int
                                hi[])
C++       void_GA::GlobalArray::distribution(int iproc, int lo[],
                                int hi[])

```

finds out the range of the global array `g_a` that process `iproc` owns. `iproc` can be any valid process ID.

The function

```

n-DFortran  subroutine nga_access(g_a, lo, hi, index, ld)
2-DFortran  subroutine ga_access(g_a, ilo, ihi, jlo, jhi, index, ld)
C           void NGA_Access(int g_a, int lo[], int hi[], void *ptr,
                           int ld[])
C++        void GA::GlobalArray::access(int lo[], int hi[], void
                           *ptr, int ld[])

```

provides access to local data in the specified patch of the array owned by the calling process. The C interface gives the pointer to the patch. The Fortran interface gives the patch address as the index (distance) from the reference address (the appropriate MA base addressing array).

The function

```

n-DFortran  subroutine nga_release(g_a, lo, hi)
2-DFortran  subroutine ga_release(g_a, ilo, ihi, jlo, jhi)
C           void NGA_Release(int g_a, lo[], int hi[])
C++        void GA::GlobalArray::release(lo[], int hi[])

```

and

```

n-DFortran  subroutine nga_release update(g_a, lo, hi)
2-DFortran  subroutine ga_release update(g_a, ilo, ihi, jlo, jhi)
C           void NGA_Release update(int g_a, int lo[], int hi[])
C++        void GA::GlobalArray::releaseUpdate(int lo[], int hi[])

```

releases access to a global array. The former set is used when the data was read only and the latter set is used when the data was accessed for writing.

Global Arrays also provide a function to compare distributions of two arrays. It is

```

Fortran     subroutine ga_compare distr(g_a, g_b)
C           void NGA_Compare distr(int g_a, int g_b)
C++        void GA::GlobalArray::compareDistr(const GA::GlobalArray *
                           g_a)

```

The only method currently available for accessing the ghost cell data for global arrays that have ghost cell data is to use the `nga_access_ghosts` function. This function is similar to the `nga_access` function already described, except that it returns an index (pointer) to the origin of the locally held patch of global array data. This local patch includes the ghost cells so the index (pointer) will be pointing to a ghost cell. The `nga_access_ghosts` function also returns the physical dimensions of the local data patch, which includes the additional ghost cells, so it is possible to access both the visible data of the global array and the ghost cells using this information. The `nga_access_ghosts` functions have the format

```

n-d Fortran  subroutine nga_access_ghosts(g_a, dims, index, ld)
C           void NGA_access_ghosts (int g_a, int dims[], void *ptr,
                           int ld[])
C++        void GA::GlobalArray::accessGhosts(int dims[], void *ptr,
                           int ld[])

```

The array `dims` comes back with the dimensions of the local data patch, including the ghost cells, for each dimension of the global array, `ptr` is an index (pointer) identifying the beginning of the local data patch, and `ld` is any array of leading dimensions for the local data patch, which also includes the ghost cells. The array `ld` is actually redundant since the information in `ld` is also contained in `dims`, but is included to maintain continuity with other GA functions.

7.1.1 Process Information

When developing a program, one needs to use characteristics of its parallel environment: process ID, how many processes are working together and what their IDs are, and what the topology of processes look like. To answer these questions, the following functions can be used.

The function

```
Fortran  integer function ga_nodeid()
C        int GA_Nodeid()
C++      int GA::GAServices::nodeid()
```

returns the GA process ID of the current process, and the function

```
Fortran  integer function ga_nnodes()
C        int GA_Nnodes()
C++      int GA::GAServices::nodes()
```


tells the number of computing processes.

The function

```
Fortran  subroutine ga_proc_topology(ga, proc, prow, pcol)
C        void NGA_Proc_topology(int g_a, int proc, int coordinates)
C++      void GA::GlobalArray::procTopology(int proc, int coordinates)
```

determines the coordinates of the specified processor in the virtual processor grid corresponding to the distribution of array `g_a`.

Example: An global array is distributed on 9 processors. The processors are numbered from 0 to 8 as shown in the following figure. If one wants to find out the coordinates of processor 7 in the virtual processor grid, by calling the function `ga_proc_topology`, the coordinates of (2,1) will be returned.



7.1.2 Cluster Information

The following functions can be used to obtain information like number of nodes that the program is running on, node ID of the process, and other cluster information as discussed below:

The function

```

Fortran  integer function _ga_cluster_nnodes()
C        int _GA_Cluster_nnodes()
C++     int _GA::GAServices::clusterNnodes()

```

returns the total number of nodes that the program is running on. On SMP architectures, this will be less than or equal to the total number of processors.

The function

```

Fortran  integer function _ga_cluster_nodeid()
C        int _GA_Cluster_nodeid()
C++     int _GA::GAServices::clusterNodeid()

```

returns the node ID of the process. On SMP architectures with more than one processor per node, several processes may return the same node id.

The function

```

Fortran  integer function _ga_cluster_nprocs(inode)
C        int _GA_Cluster_nprocs(int inode)
C++     int _GA::GAServices::clusterNprocs(int inode)

```

returns the number of processors available on node inode.

The function


```

Fortran  integer function _ga_cluster_procid(inode, iproc)
C        int _GA_Cluster_procid(int inode, int iproc)
C++     int _GA::GAServices::clusterProcid(int inode, int iproc)

```

returns the processor id associated with node inode and the local processor id iproc. If node inode has N processors, then the value of iproc lies between 0 and N-1.

Example: 2 nodes with 4 processors each. Say, there are 7 processes created. Assume 4 processes on node 0 and 3 processes on node 1. In this case: number of nodes=2, node id is either 0 or 1 (for example, nodeid of process 2 is 0), number of processes in node 0 is 4 and node 1 is 3. The global rank of each process is shown in the figure and also the local rank (rank of the process within the node.i.e.cluster_procid) is shown in the paranthesis.



7.2 Memory Availability

Even though the memory management does not have to be performed directly by the user, Global Arrays provide functions to verify the memory availability. Global Arrays provide the following information:

1. How much memory has been used by the allocated global arrays.
2. How much memory is left for allocation of new the global arrays.
3. Whether the memory in global arrays comes from the Memory Allocator (MA).
4. Is there any limitation for the memory usage by the Global Arrays.

The function

```

Fortran integer function ga_inquire_memory()
C       size_t GA_Inquire_memory()
C++    size_t GA::GAServices::inquireMemory()

```

answers the first question. It returns the amount of memory (in bytes) used in the allocated global arrays on the calling processor.

The function

```

Fortran integer function ga_memory_avail()
C       size_t GA_Memory_avail()
C++    size_t GA::GAServices::memoryAvailable()

```

answers the second question. It returns the amount of memory (in bytes) left for allocation of new global arrays on the calling processor.

Memory Allocator(MA) is a library of routines that comprises a dynamic memory allocator for use by C, Fortran, or mixed-language applications. Fortran-77 applications require such a library because the language does not support dynamic memory allocation. C (and Fortran-90) applications can benefit from using MA instead of the ordinary malloc() and free() routines because of the extra features MA provides. The function

```

Fortran logical function ga_uses_ma()
C       int GA_Uses_ma()

```

```
C++      int GA::GAServices::usesMA()
```

tells whether the memory in Global Arrays comes from the Memory Allocator (MA) or not.

The function

```
Fortran logical function_ga_memory_limited()  
C      int GA_Memory_limited()  
C++    int GA::GAServices::memoryLimited()
```

Indicates if a limit is set on memory usage in Global Arrays on the calling processor.

7.3 Message–Passing Wrappers to Reduce/Broadcast Operations

Global Arrays provide convenient operations for broadcast/reduce regardless of the message–passing library the process is running with.

The function

```
Fortran subroutine ga_brdcst(type, buf, lenbuf, root)  
C      void GA_Brdcst(void *buf, int lenbuf, int root)  
C++    void GA::GAServices::brdcst(void *buf, int lenbuf, int root)
```

broadcasts from process root to all other processes a message buffer of length lenbuf.

The functions

```
Fortran subroutine ga_igop(type, x, n, op)  
          subroutine ga_dgop(type, x, n, op)  
C      void GA_Igop(long x[], int n, char *op)  
          void GA_Dgop(double x[], int n, char *op)  
C++    void GA::GAServices::igop(long x[], int n, char *op)  
          void GA::GAServices::dgop(double x[], int n, char *op)
```

'sum' elements of $X(I:N)$ (a vector present on each process) across all nodes using the communicative operator op, The result is broadcasted to all nodes. Supported operations include

+, *, Max, min, Absmax, absmin

The integer version also includes the **bitwise OR** operation.

These operations unlike `ga_sync`, do not include embedded `ga_genceoperations`.

7.4 Others

There are some other useful functions in Global Arrays. One group is about inquiring the array attributes. Another group is about printing the array or part of the array.

7.4.1 Inquire

A global array is represented by a handle. Given a handle, one can get the array information, such as the array name, memory used, array data type, and array dimension information, with the help of following functions.

The functions

```

n-D Fortran  subroutine nga_inquire(g_a, type, ndim, dims)
2-D Fortran  subroutine nga_inquire(g_a, type, dim1, dim2)
C           void NGA_Inquire(int g_a, int *type, int *ndim, int
                        dims[])
C++         void GA::GlobalArray::inquire(int *type, int *ndim, int
                        dims[])

```

return the data type of the array, and also the dimensions of the array.

The function

```

Fortran  subroutine ga_inquire_name(g_a, array_name)
C       char* GA_Inquire_name(int g_a)
C++     char* GA::GlobalArray::inquireName()

```

finds out the name of the array.

One can also inquire the memory being used with `ga_inquire_memory`(discussed above).

7.4.2 Print

Global arrays provide functions to print

1. content of the global array
2. content of a patch of global array
3. the status of array operations
4. a summary of allocated arrays

The function

```

Fortran  subroutine ga_print(g_a)
C       void GA_Print(int g_a)
C++     void GA::GlobalArray::print()

```

prints the entire array to the standard output. The output is formatted.

A utility function is provided to print data in the patch, which is

```

Fortran  subroutine nga_print_patch(g_a, lo, hi, pretty)
C       void NGA_Print_patch(int g_a, int lo[], int hi[], int pretty)
C++     void GA::GlobalArray::printPatch(int lo[], int hi[], int
                        pretty)

```

One can either specify a formatted output (set `pretty` to one) where the output is formatted and rows/columns are labeled, or (set `pretty` to zero) just dump all the elements of this patch to the standard output without any formatting.

The function

```
Fortran  subroutine ga_print_stats()
C        void GA_Print_stats()
C++      void GA::GAServices::printStats()
```

prints the global statistics information about array operations for the calling process, including

- number of calls to the GA create/duplicate, destroy, get, put, scatter, gather, and read_and_inc operations
- total amount of data moved in the GA primitive operations
- amount of data moved in GA primitive operations to logically remote locations
- maximum memory consumption in global arrays, the "high-water mark"

The function

```
Fortran  subroutine ga_print_distribution(g_a)
C        void GA_Print_distribution(int g_a)
C        void GA::GlobalArray::printDistribution()
```

prints the global array distribution. It shows mapping array data to the processes.

The function

```
Fortran  subroutine ga_summarize(verbose)
C        void GA_Summarize(int verbose)
C++      void GA::GAServices::summarize(int verbose)
```

prints info about allocated arrays. `verbose` can be either one or zero.

7.4.3 Miscellaneous

The function

```
Fortran  subroutine ga_check_handle(g_a, string)
C        void GA_Check_handle(int g_a, char *string)
C++      void GA::GlobalArray::checkHandle(char *string)
```

checks if the global array handle `g_a` represents a valid array. The `string` is the message to be printed when the handle is invalid.

GA++: C++ Bindings for Global Arrays

8.1 Overview

GA++ provides a C++ interface to global arrays (GA) libraries. Here is the doxygen documentation of GA++: <http://www.emsl.pnl.gov:2080/docs/global/ga++/index.html> The GA C++ bindings are a layer built directly on top of the GA C bindings. GA++ provides new names for the C bindings of GA functions (For example, `GA_Add_patch()` is renamed as `addPatch()`).

8.2 GA++ Classes

All GA classes (`GAServices`, `GlobalArray`) are declared within the scope of GA namespace.

Namespace issue: Although namespace is part of ANSI C++ standard, not all C++ compilers support namespaces (A non-instantiable GA class is provided for implementations using compilers without namespace).

Note: define the variable `_GA_USENAMESPACE_` as 0 in `ga++.h` if your compiler doesnot support namespaces.

```
namespace GA {
    class GAServices;
    class GlobalArray;
};
```

Current implementation has no derived classes (no (virtual) inheritance), templates or exception handling. Eventually, more object oriented functionalities will be added, and standard library facilities will be used without affecting the performance.

8.3 Initialization and Termination:

GA namespace has the following static functions for initialization and termination of Global Arrays.

GA::Initialize():

Initialize Global Arrays, allocates and initializes internal data structures in Global Arrays. This is a collective operation.

GA::Terminate():

Delete all active arrays and destroy internal data structures. This is a collective operation.

```
namespace GA {
    _GA_STATIC_ void Initialize(int argc, char *argv[], size_t limit = 0);
    _GA_STATIC_ void Initialize(int argc, char *argv[], unsigned long heapSize, unsigned long
stackSize, int type, size_t limit = 0);
    _GA_STATIC_ void Terminate();
};
```

Example:

```
#include <iostream.h>
#include "ga++.h"

int
main(int argc, char **argv) {
    GA::Initialize(argc, argv, 0);
    cout << "Hello World\n";
    GA::Terminate();
}
```

8.4 GAServices:

GAServices class has member functions that does all the global operations (non-array operations) like Process Information (number of processes, process id, ..), Inter-process Synchronization (sync, lock, broadcast, reduce,..), etc.,

SERVICES Object:

GA namespace has a global "SERVICES" object (of type "GAServices"), which can be used to invoke the non-array operations. To call the functions (for example, sync()), we invoke them on this SERVICES object (for example, GA::SERVICES.sync()). As this object is in the global address space, the functions can be invoked from anywhere inside the program (provided the ga++.h is included in that file/program).

8.5 Global Array:

GlobalArray class has member functions that perform:

- Array operations
 - One-sided (get/put),
 - Collective array operations,
 - Utility operations, etc.,
-